

Table 1: Disk Parameters

disk pack configuration		2 double-sided disks
read/write heads		4
blocks per track		400
tracks per surface		20,000
block size	B	512 bytes
rotational velocity	p	7200 rpm
average seek time	s	6 msec
block pointer size	P	6 bytes
record pointer size	P _R	7 bytes

Table 2: EMPLOYEE Relation

number of records	500,000		
storage	unspanned, contiguous, ordered on PK (EMPID)		
schema	attribute	domain	size (bytes)
	FNAME	CHAR(15)	20
	LNAME	CHAR(15)	20
	<u>EMPID</u>	CHAR(9)	9
	BDATE	DATE	4
	ADDRESS	CHAR(30)	30
	SEX	CHAR	1
	SALARY	DECIMAL(10,2)	8
	SUPERSSN	CHAR(9)	9
	DNO	INT	4
	TITLE	CHAR(15)	15

Record Storage: Answer the following for the EMPLOYEE relation:

- (a) Compute the record size (R):
- (b) Compute the blocking factor (bfr):
- (c) How many blocks are required to store the relation?

Primary Indexes: Compute the following for a *primary index* on the EMPLOYEE relation.

- (a) What is the index record size?
- (b) What is the index blocking factor (bfr_i)?
- (c) How many disk blocks are required to store a *single-level* primary index?
- (d) How many disk blocks are required to store a *multi-level* primary index?
(Assume the top level is one block.)
- (e) How many block reads are required to execute the following parameterized query:
`SELECT * FROM EMPLOYEE WHERE EMPID = [parameter];`
With no index?
With a single-level primary index?
With a multiple-level primary index?
- (f) Estimate the *time* required to execute the query in (e), for each of the three conditions.

Clustering Indexes: Suppose the EMPLOYEE relation is ordered on field DNO (department number), rather than the primary key. There are 100 departments and employees are evenly distributed among the departments.

- (a) What is the index blocking factor (bfr_i)?
- (b) How many blocks are required to store a clustering index on DNO?
- (c) Would a multi-level index on DNO be effective?

Secondary Indexes:

- (a) Assume the company is 60% male and 40% female. Would a secondary index on the field SEX be effective?
- (b) Assume there are 2000 departments and the employees are evenly distributed among those departments. Assume that the primary key EMPID orders the EMPLOYEE relation, and that a multi-level primary index exists.
Design a secondary index on DNO.
- (c) How many blocks are required to store your secondary index?
- (d) How many block accesses (on average) are required to execute the following parameterized query, using your secondary index?
`SELECT * FROM EMPLOYEE WHERE DNO = [parameter];`
- (e) Design a secondary index on employee names.