

Advanced Database Management Systems

**Lecture 8 – Sections 8.4-8.5
SQL Queries**

SQL

- **SQL** (Structured Query Language)
is the standard language for commercial DBMSs
- **SEQUEL** (Structured English QUery Language)
was originally defined by IBM for **SYSTEM R**
 - mid 1970s
 - unofficial pronunciation (**see**-kwuhl) still sometimes used
- standardization of SQL began in the 80s
- current standard is SQL-99
 - subsequent revisions are not fully accepted by all vendors
- SQL is more than a query language:
it includes a DDL, DML and admin commands

SQL commands

- Administration:
 - CREATE DATABASE
 - CREATE SCHEMA
 - SET ROLE
 - GRANT PRIVILEGES
- Data Definition:
 - CREATE TABLE
 - ALTER TABLE
 - DROP TABLE
 - CREATE VIEW
- Data Modification:
 - INSERT
 - DELETE
 - UPDATE
- Queries:
 - SELECT

48 commands listed in
SQL in a Nutshell

SQL Queries

- Queries in SQL are variations of the SELECT command
- Basic SQL queries correspond to the following relational algebra operations:
 - select σ
 - project π
 - cross product \times
 - joins must be expressed as σ and \times

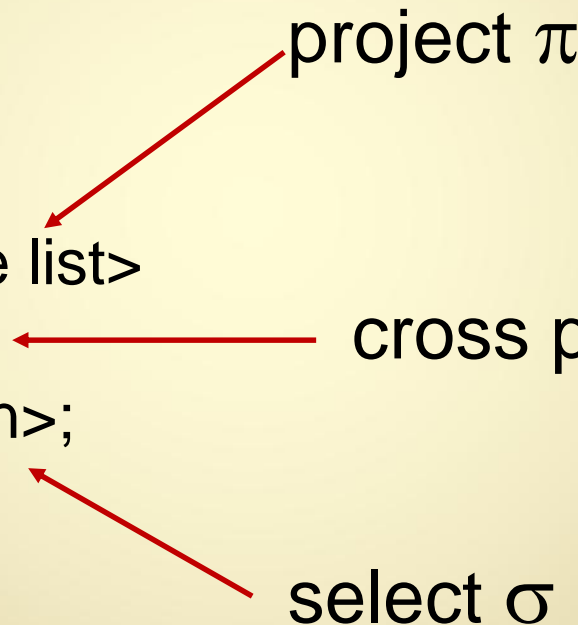
Basic SELECT Command

SELECT <attribute list>
FROM <table list>
WHERE <condition>;

project π

cross product \times

select σ



Single Table Queries (σ and π)

$\pi_{Ssn}(\sigma_{Salary > 60000}(\text{EMPLOYEE}))$

```
SELECT Ssn  
FROM EMPLOYEE  
WHERE Salary > 60000;
```

$\pi_{City, State}(\sigma_{Airport_code = 'SFO'}(\text{AIRPORT}))$

```
SELECT City, State  
FROM AIRPORT  
WHERE Airport_code = 'SFO';
```

Join as Select & Cross

- In the basic SELECT/FROM/WHERE form, joins must be expressed as using σ and \times

$$\pi_{Lname, Dname} (EMPLOYEE \bowtie_{Ssn=Mgr_ssn} DEPARTMENT)$$

$$\pi_{Lname, Dname} (\sigma_{Ssn=Mgr_ssn} (EMPLOYEE \times DEPARTMENT))$$

```
SELECT Lname, Dname  
FROM EMPLOYEE, DEPARTMENT  
WHERE Ssn = Mgr_ssn;
```

Basic SQL Queries

- Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT    FNAME, LNAME, ADDRESS
FROM      EMPLOYEE, DEPARTMENT
WHERE     DNAME='Research' AND DNUMBER=DNO
```

selection
condition

join
condition

Basic SQL Queries

- For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
SELECT PNUMBER, DNUM, LNAME, BDATE, ADDRESS
```

```
FROM PROJECT, DEPARTMENT, EMPLOYEE
```

```
WHERE DNUM=DNUMBER AND MGRSSN=SSN
```

```
AND PLOCATION='Stafford'
```

join
PROJECT and
DEPARTMENT



select



join
DEPARTMENT
and EMPLOYEE



Tuple Variables (Aliases)

- We can give names to the tuples coming from each of the input relations

```
SELECT      E.Lname, D.Dname  
FROM        EMPLOYEE E, DEPARTMENT D  
WHERE       E.Ssn = D.Mgr_ssn;
```

- This can disambiguate common attribute names and improve readability

Renaming Attributes

- Attributes can also be renamed in the FROM clause
 - similar to alternate rename syntax in the algebra

SELECT	Fn, Ln
FROM	EMPLOYEE E(Fn, Mi, Ln, Bd, Ad, Sx, Sl, Sssn, Dn)
WHERE	Dn = 4;

Self Join

- For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

```
SELECT      E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM        EMPLOYEE E, EMPLOYEE S
WHERE       E.SUPERSSN=S.SSN
```

- Aliases are necessary for this query
- Think of E and S as two different copies of EMPLOYEE
 - E represents employees in role of *supervisees* and
S represents employees in role of *supervisors*

Aliases: alternate syntax

- Can also use the AS keyword to specify aliases

```
SELECT      E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM        EMPLOYEE AS E, EMPLOYEE AS S
WHERE       E.SUPERSSN=S.SSN
```

- Can also simply use the relation names
(when non-ambiguous)

```
SELECT      EMPLOYEE.Lname, DEPARTMENT.Dname
FROM        EMPLOYEE, DEPARTMENT
WHERE       EMPLOYEE.Ssn = DEPARTMENT.Mgr_ssn;
```

No $\sigma \rightarrow$ No WHERE

- If there are no selection (or join) conditions, the WHERE clause can be omitted

SELECT Ssn
FROM EMPLOYEE π_{Ssn} EMPLOYEE

- Two or more relations in FROM clause with no join is a *CROSS PRODUCT*

SELECT Lname, Dname
FROM EMPLOYEE, DEPARTMENT

$\pi_{Lname, Dname} (EMPLOYEE \times DEPARTMENT)$

No $\pi \rightarrow *$

- To retrieve all the attribute values of the selected tuples, use *, which stands for *all the attributes*

```
SELECT      *  
FROM        EMPLOYEE  
WHERE       DNO=5
```

```
SELECT      *  
FROM        EMPLOYEE, DEPARTMENT  
WHERE       DNAME='Research' AND  
            DNO=DNUMBER
```

Tables as Sets → DISTINCT

- SQL does not treat a relation as a set; duplicate tuples can appear
- To eliminate duplicate tuples in a query result, the keyword **DISTINCT** is used

```
SELECT    SALARY  
FROM      EMPLOYEE
```

may contain duplicates

```
SELECT    DISTINCT SALARY  
FROM      EMPLOYEE
```

duplicates eliminated



Set Operations

- union operation (UNION)
intersection (INTERSECT)
set difference (MINUS, sometimes called EXCEPT)
 - some implementations of SQL do not support all set operations
- Set operation results **are sets** of tuples
duplicate tuples are eliminated from the result
- The set operations apply only to *union compatible relations*:
the two relations must have the same attributes and
the attributes must appear in the same order

Set Operations: Example

- List project numbers for all projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT      PNAME
FROM          PROJECT, DEPARTMENT, EMPLOYEE
WHERE        DNUM=DNUMBER AND MGRSSN=SSN AND LNAME='Smith')
```

UNION

```
(SELECT      PNAME
FROM          PROJECT, WORKS_ON, EMPLOYEE
WHERE        PNUMBER=PNO AND ESSN=SSN AND NAME='Smith')
```

Multiset Operations

- UNION ALL, INTERSECT ALL, EXCEPT ALL
- Multiset operation results are multisets of tuples
duplicate tuples are not eliminated

```
(SELECT      PNAME
FROM        PROJECT, DEPARTMENT, EMPLOYEE
WHERE       DNUM=DNUMBER AND MGRSSN=SSN AND LNAME='Smith')
      UNION ALL
(SELECT      PNAME
FROM        PROJECT, WORKS_ON, EMPLOYEE
WHERE       PNUMBER=PNO AND ESSN=SSN AND NAME='Smith')
```

WHERE Clause

- WHERE clause is a general boolean expression
- Boolean operators:
AND, OR, NOT
- Comparison operators:
=, <, <=, >, >=, <>
- String comparison operators:
LIKE
- Parentheses can be used to set precedence
- String literals can be enclosed in "... " or '...'

String Comparison

- The **LIKE** comparison operator is used to compare partial strings

- Two wildcard characters are used:

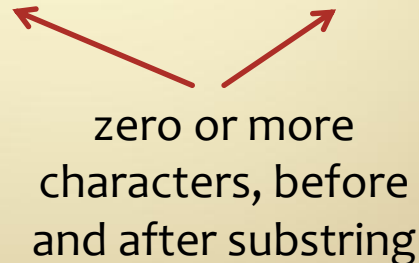
'%' replaces an arbitrary number of characters

'_' replaces a single arbitrary character

String Comparison Example

- Retrieve all employees whose address is in Houston, Texas.
- The value of the ADDRESS attribute must contain the substring “Houston, TX”.

```
SELECT    FNAME, LNAME  
FROM      EMPLOYEE  
WHERE     ADDRESS LIKE '%Houston, TX%'
```



zero or more
characters, before
and after substring

String Comparison Example

- Retrieve all employees who were born during the 1960s.
 - '6' must be the 3rd character of the 10 character date string

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       BDATE LIKE  '__ 6 _____'
```

- Following would also work:

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       BDATE LIKE  '__ 6 %'
```

assumes date format is YYYY-MM-DD

Arithmetic Operation

- The standard arithmetic operators '+', '-', '*', and '/' can be applied to numeric values in an SQL query result
- Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
SELECT      FNAME, LNAME, 1.1*SALARY
FROM        EMPLOYEE, WORKS_ON, PROJECT
WHERE       SSN=ESSN AND PNO=PNUMBER
            AND PNAME='ProductX'
```


Aggregate Functions

- Aggregate functions are applied to result attributes **COUNT, SUM, MAX, MIN, and AVG**
- Find the maximum salary, the minimum salary, and the average salary among all employees.

```
SELECT      MAX(Salary), MIN(Salary), AVG(Salary)
FROM        EMPLOYEE
```

- Find the total salary paid to employees who work for the 'Research' department.

```
SELECT      SUM(Salary)
FROM        EMPLOYEE, DEPARTMENT
WHERE       Dno=Dnumber AND Dname='Research'
```

Aggregate Functions

- Retrieve the total number of employees in the company and the number of employees in the Research' department.

```
SELECT    COUNT (*)  
FROM      EMPLOYEE
```

```
SELECT    COUNT (*)  
FROM      EMPLOYEE, DEPARTMENT  
WHERE     DNO=DNUMBER AND DNAME='Research'
```

Join as X and σ

```
mysql> SELECT * FROM r;
```

x	y
3	4
5	6
7	8
9	6

```
mysql> SELECT * FROM s;
```

a	b
2	3
4	7

```
mysql> SELECT * FROM r,s WHERE y=a;
```

x	y	a	b
3	4	4	7

$$R \bowtie_{y=a} S$$

$$\sigma_{y=a} (R \times S)$$

Explicit Join

- Joins can be explicitly stated in the FROM clause.

```
SELECT * FROM (r JOIN s ON y=a) ;
```

x	y	a	b
3	4	4	7

$R \bowtie_{y=a} S$

Left/Right Outer Join

```
SELECT * FROM (r LEFT JOIN s ON y=a) ;
```

x	y	a	b
3	4	4	7
5	6	NULL	NULL
7	8	NULL	NULL
9	6	NULL	NULL

```
SELECT * FROM (r RIGHT JOIN s ON y=a) ;
```

x	y	a	b
NULL	NULL	2	3
3	4	4	7

Full Outer Join

```
SELECT * FROM r FULL OUTER JOIN s ON y=a;
```

mysql doesn't support full outer join,
so we'll substitute an equivalent query:

```
(SELECT * FROM r LEFT JOIN s ON y=a) UNION  
(SELECT * FROM r RIGHT JOIN s ON y=a);
```

x	y	a	b
3	4	4	7
5	6	NULL	NULL
7	8	NULL	NULL
9	6	NULL	NULL
NULL	NULL	2	3

Ordering Results

- An ORDER BY clause can be added to order the result tuples

<code>SELECT * FROM t;</code>	<code>SELECT * FROM t ORDER BY j;</code>	<code>SELECT * FROM t ORDER BY i;</code>
<code>+-----+-----+</code>	<code>+-----+-----+</code>	<code>+-----+-----+</code>
<code> i j </code>	<code> i j </code>	<code> i j </code>
<code>+-----+-----+</code>	<code>+-----+-----+</code>	<code>+-----+-----+</code>
<code> 10 ten </code>	<code> 11 eleven </code>	<code> 4 four </code>
<code> 11 eleven </code>	<code> 4 four </code>	<code> 10 ten </code>
<code> 20 twenty </code>	<code> 10 ten </code>	<code> 11 eleven </code>
<code> 4 four </code>	<code> 20 twenty </code>	<code> 20 twenty </code>
<code>+-----+-----+</code>	<code>+-----+-----+</code>	<code>+-----+-----+</code>

ORDER BY Examples

- order by Lname first,
then by Fname if Lname is the same:

```
SELECT Lname, Fname  
FROM Employee  
WHERE salary > 60000  
ORDER BY Lname, Fname
```

- order by Lname in ascending order,
then by salary in descending order

```
SELECT Lname, salary  
FROM Employee  
WHERE salary > 60000  
ORDER BY Lname ASC, salary DESC
```


Grouping

- Forms groups (subsets) of result tuples before applying aggregate functions
- Example: count the number of employees in each department
(group employees by DNO, then count tuples in each group)

```
SELECT Dno, COUNT(*)  
FROM Employee  
GROUP BY Dno
```

$\text{DNO } \mathcal{F}_{\text{COUNT}} * (\text{EMPLOYEE})$

+-----+-----+		
Dno	COUNT	
+-----+-----+		
8	120	
22	238	
7	82	
20	169	
+-----+-----+		

GROUP BY Example

- For each project, get the project name, project number and the number of employees working on that project

```
SELECT Pnumber, Pname, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE Pnumber = Pno  
GROUP BY Pnumber, Pname
```

Attributes in SELECT clause must be aggregates
or must appear in the GROUP BY clause

Filtering Groups: HAVING

- We can throw away some groups by adding a condition in a HAVING clause
- example:
for each project *that has more than two employees*,
get the project name, project number and
the number of employees working on that project

```
SELECT Pnumber, Pname, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE Pnumber = Pno  
GROUP BY Pnumber, Pname  
HAVING COUNT(*) > 2
```

GROUP BY Examples

```
SELECT * FROM e;
```

eid	salary	dept
E01	65000	ADMIN
E12	58400	ENGR
E08	76900	ENGR
E23	63800	ADMIN
E07	56900	ADMIN
E27	76400	ENGR
E14	48000	TEST

```
SELECT COUNT(*) FROM e
GROUP BY dept;
```

count(*)
3
3
1

```
SELECT dept, COUNT(*)
FROM e GROUP BY dept;
```

dept	count(*)
ADMIN	3
ENGR	3
TEST	1

GROUP BY Examples

```
SELECT dept, COUNT(*)  
FROM e  
GROUP BY dept  
HAVING COUNT(*) > 1;
```

dept	count(*)
ADMIN	3
ENGR	3

```
SELECT dept, AVG(salary)  
FROM e  
GROUP BY dept  
HAVING COUNT(*) > 1;
```

dept	AVG(salary)
ADMIN	61900
ENGR	70566.66667

Nested Queries

- Nested queries can be used as set values in the WHERE clause
- Set comparison operators
 - IN – set membership (“is in”, \in)
 - EXISTS – set not empty (\exists)
 - ALL – applies to all set members (\forall)
 - ANY – applies to any set member
 - CONTAINS – proper superset

Nested Queries

- find all employees who work on a project with John Smith

```
SELECT Lname, Fname
FROM EMPLOYEE E1, WORKS_ON W1
WHERE E1.SSN = W1.ESSN
      AND W1.Pno IN (SELECT Pno
                      FROM EMPLOYEE E2, WORKS_ON W2
                      WHERE E2.SSN = W2.ESSN
                         AND E2.Fname = "John"
                         AND E2.Lname = "Smith")
```

Nested Queries

- find the highest paid employee in department 5

```
SELECT Lname, Fname
FROM EMPLOYEE E1
WHERE E1.Dno=5
      AND E1.Salary > ALL (SELECT E2.Salary
                           FROM EMPLOYEE E2
                           WHERE E2.Dno=5)
```


Nested Queries

- List names of managers who have dependents

```
SELECT Lname, Fname
FROM EMPLOYEE E1
WHERE EXISTS (SELECT *
              FROM DEPENDENT D1
              WHERE E1.Ssn = D1.Essn)
AND
EXISTS (SELECT *
       FROM DEPARTMENT D2
       WHERE E1.Ssn = D2.Mgr_ssn)
```

This is an example of a *correlated nested query*,
since the nested queries refer to the relations in the outer query.

Nested Queries


- List names of employees who work on all projects controlled by department 5

```
SELECT Lname, Fname
FROM EMPLOYEE E
WHERE (SELECT W.Pno
        FROM WORKS_ON W
        WHERE E.Ssn = W.Essn)
CONTAINS
(SELECT P.Pnumber
    FROM PROJECT P
    WHERE P.Dnum=5)
```

Nested Queries

- List names of all projects controlled by department 5 or department 7

```
SELECT P.Pname  
FROM PROJECT P  
WHERE P.Dnum IN (5,7)
```



explicit set of values

SELECT: Syntax Summary

SELECT <attribute and function list>

FROM <table list>

WHERE <condition>

GROUP BY <grouping attributes>

HAVING <group condition>

ORDER BY <attribute list>

required

optional

SELECT: conceptual execution

1. FROM: cross product of tables
2. WHERE: select tuples
3. GROUP BY: group tuples
4. HAVING: filter groups
5. SELECT: project attributes and apply aggregates
6. ORDER BY: sort the tuples

This is not an efficient way to execute the query,
simply a way to define the meaning of the query conceptually.

EXERCISE 1: Queries

1. First and last name of employees who have no supervisor.
2. First and last name of employees supervised by Franklin Wong.
3. Last name of employees who have dependents.
4. Last name of employees who have daughters.
5. Last name of employees in department 5 who work more than 10 hours/week on ProductX.
6. Last name of supervisors of employees in department 5 who work more than 10 hours/week on ProductX.
7. First and last names of all department managers.
8. Salaries of all employees who have worked on the Reorganization project.
9. SSN of all employees who have worked on a project that is controlled by a department different than the department that they are assigned to.
10. Last name of all employees who are not married.

EXERCISE 2: Queries

1. List all airplane types that can land at any airport in San Francisco.
2. List the ids and number of seats for all airplanes that can land at any airport in Chicago.
3. List the name and phone number of all customers with a seat reserved on a flight that leaves Chicago O'Hara airport (ORD) on October 31, 2008.
4. List all airlines that have seats available for flights leaving Los Angeles (LAX) on September 25, 2008.
5. List all airlines that operate at San Jose International Airport (SJC).

EXERCISE 3: Queries

1. Count the number of overdue books.
2. How many books by author Harry Crews are in the database?
3. Determine the number of library cards assigned to each borrower phone number.
4. Find names of all borrowers who do not have any book loans.
5. Do any library branches have every book?

EXERCISE 1: Schema

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

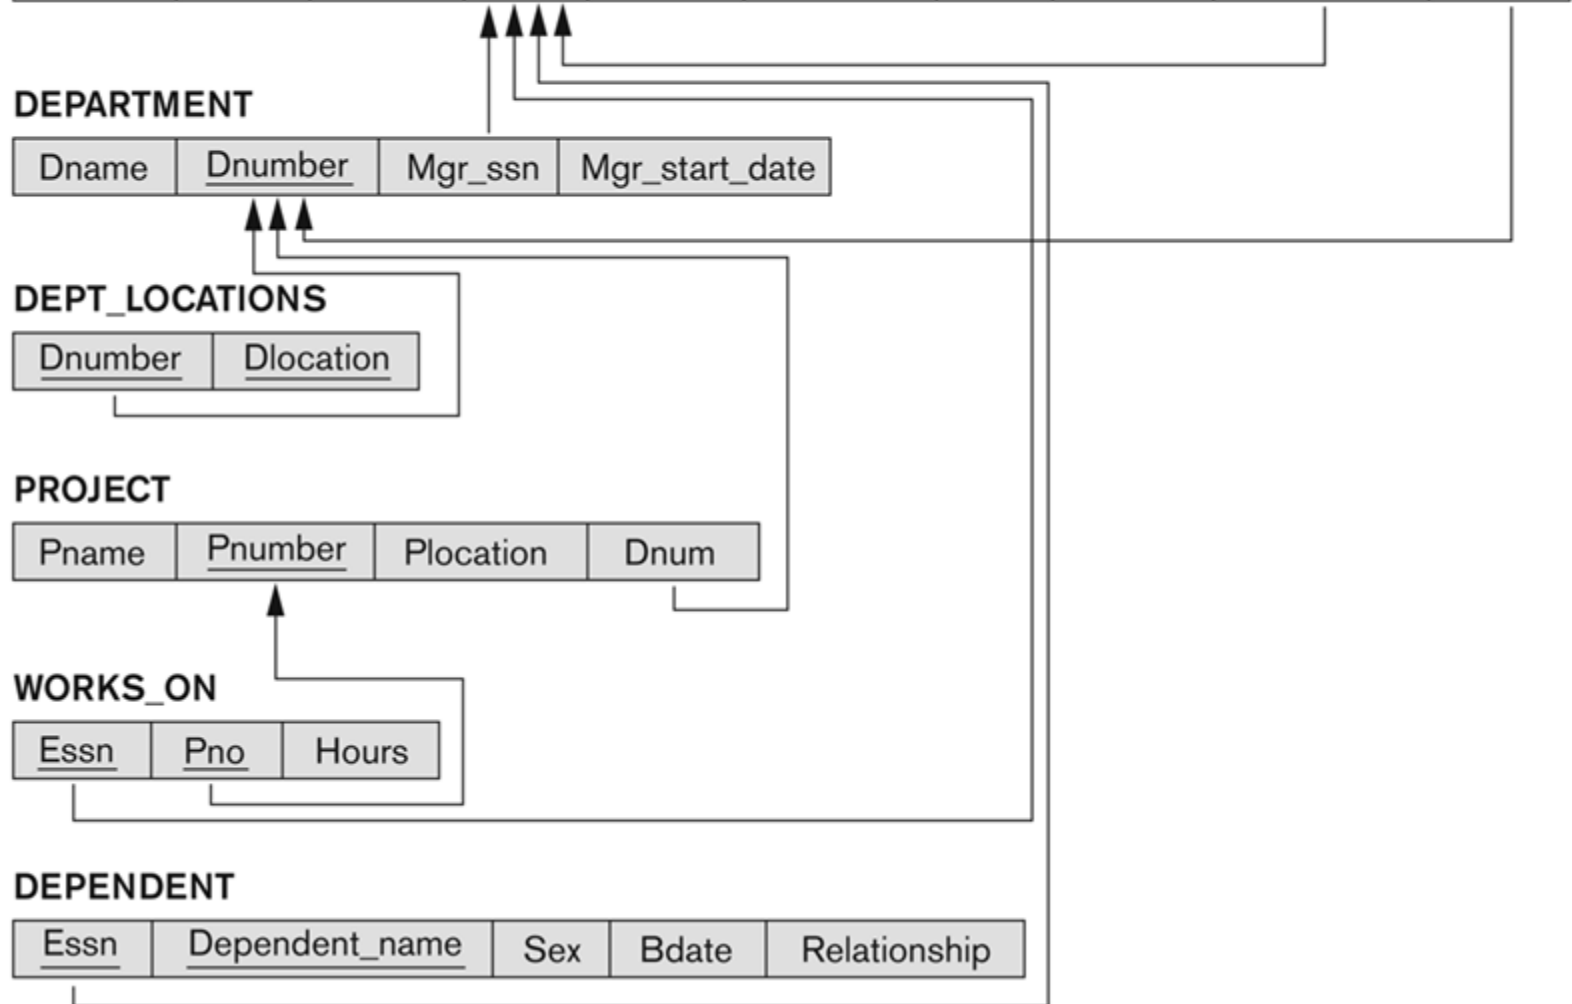
Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------



EXERCISE 1:

Instance

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

Exercise 2: Schema

AIRPORT

<u>Airport_code</u>	Name	City	State
---------------------	------	------	-------

FLIGHT

<u>Flight_number</u>	Airline	Weekdays
----------------------	---------	----------

FLIGHT_LEG

<u>Flight_number</u>	<u>Leg_number</u>	Departure_airport_code	Scheduled_departure_time
		Arrival_airport_code	Scheduled_arrival_time

LEG_INSTANCE

<u>Flight_number</u>	<u>Leg_number</u>	<u>Date</u>	Number_of_available_seats	Airplane_id
	Departure_airport_code	Departure_time	Arrival_airport_code	Arrival_time

FARE

<u>Flight_number</u>	<u>Fare_code</u>	Amount	Restrictions
----------------------	------------------	--------	--------------

AIRPLANE_TYPE

<u>Airplane_type_name</u>	Max_seats	Company
---------------------------	-----------	---------

CAN_LAND

<u>Airplane_type_name</u>	<u>Airport_code</u>
---------------------------	---------------------

AIRPLANE

<u>Airplane_id</u>	Total_number_of_seats	Airplane_type
--------------------	-----------------------	---------------

SEAT_RESERVATION

<u>Flight_number</u>	<u>Leg_number</u>	<u>Date</u>	<u>Seat_number</u>	Customer_name	Customer_phone
----------------------	-------------------	-------------	--------------------	---------------	----------------

EXERCISE 3:

Schema

BOOK

<u>Book_id</u>	Title	Publisher_name
----------------	-------	----------------

BOOK_AUTHORS

<u>Book_id</u>	<u>Author_name</u>
----------------	--------------------

PUBLISHER

<u>Name</u>	Address	Phone
-------------	---------	-------

BOOK_COPIES

<u>Book_id</u>	<u>Branch_id</u>	No_of_copies
----------------	------------------	--------------

BOOK_LOANS

<u>Book_id</u>	<u>Branch_id</u>	<u>Card_no</u>	Date_out	Due_date
----------------	------------------	----------------	----------	----------

LIBRARY_BRANCH

<u>Branch_id</u>	Branch_name	Address
------------------	-------------	---------

BORROWER

<u>Card_no</u>	Name	Address	Phone
----------------	------	---------	-------

