

Advanced Database Management System

**Lecture 10 – Sections 8.7, 8.8, 24.1
Views, Semantic Constraints
and Triggers**

REVIEW: SQL Queries

- Count the number of men and women in each department.

Report: (department name, # of women, # of men)

EMPLOYEE(..., SSN, ..., SEX, ..., DNO)
domain(SEX) = { 'M', 'F' }

DEPARTMENT(DNAME, DNUMBER, ...)



1: Count women/men per department

R_f = select dno as dno1, count(*) as cnt_f
from employee where sex='F'
group by dno

dno1	cnt_f
4	2
5	1

dno2	cnt_m
1	1
4	1
5	3

R_m = select dno as dno2, count(*) as cnt_m
from employee where sex='M'
group by dno

2a: combine the results

+	-----	+	-----	+
	dno1		cnt_f	
+	-----	+	-----	+
	4		2	
	5		1	
+	-----	+	-----	+



dno1=dno2


+	-----	+	-----	+
	dno2		cnt_m	
+	-----	+	-----	+
	1		1	
	4		1	
	5		3	
+	-----	+	-----	+

Result:

+	-----	+	-----	+	-----	+	-----	+
	dno1		cnt_f		dno2		cnt_m	
+	-----	+	-----	+	-----	+	-----	+
	4		2		4		1	
	5		1		5		3	
+	-----	+	-----	+	-----	+	-----	+

close ... but we lost department 1!

2b: combine the results

-----+	-----+		-----+	-----+
dno1	cnt_f		dno2	cnt_m
-----+	-----+		-----+	-----+
4	2	=  =	1	1
5	1	dno1=dno2	4	1
-----+	-----+		5	3
			-----+	-----+

Result:

-----+	-----+	-----+	-----+
dno1	cnt_f	dno2	cnt_m
-----+	-----+	-----+	-----+
4	2	4	1
5	1	5	3
NULL	NULL	1	1
-----+	-----+	-----+	-----+

full outer join gives everything we need ...
unfortunately, MySql doesn't support full outer join.

2c: alternate approach

Since we'll need department names, left join the initial results with DEPARTMENT, then join those results.

$$\begin{aligned} & (\text{Department} \Rightarrow \bowtie_{\text{dnumber}=\text{dno1}} R_f) \\ & \quad \bowtie_{\text{dnumber}=\text{dnumber}} \\ & (\text{Department} \Rightarrow \bowtie_{\text{dnumber}=\text{dno1}} R_m) \end{aligned}$$

The left joins force all department numbers to be represented in the subqueries, and they will all find a match in the third join.

3: subqueries

```
(select department.dnumber as department1, cnt_f as women  
from  
department left join  
(select dno as dno1, count(*) as cnt_f  
from employee where sex='F'  
group by dno) as R1 on dno1=department.dnumber) as R3
```

Result will look like (department1, cnt_f) and will include all departments. Departments with no women will have NULL in second column.

Subquery for counting men will be similar.

4: solution in SQL

select **department1** as **department**, **women**, **men**
from

(select **department.dnumber** as **department1**, **cnt_f** as **women**
from

department left join

(select **dno** as **dno1**, count(*) as **cnt_f**
from **employee** where **sex**='F'

group by **dno**) as **R1** on **dno1=department.dnumber**) as **R3**,

(select **department.dnumber** as **department2**, **cnt_m** as **men**
from

department left join

(select **dno** as **dno1**, count(*) as **cnt_m**
from **employee** where **sex**='M'

group by **dno**) as **R2** on **dno1=department.dnumber**) as **R4**

where **department1=department2**;

Views

Views in SQL

- A **view** is a **virtual table** that is derived from other tables
 - Virtual → computed, not stored
- Allows full query operations
 - Behaves as a table when queried
- Allows for limited update operations
 - Updates must be mapped down to tables that contribute to the view

Specification of Views

- SQL command: **CREATE VIEW**
 - a table (view) name
 - a possible list of attribute names
(for example, when arithmetic operations are specified or when we want the names to be different from the attributes in the base relations)
 - a query to specify the table contents

SQL Views: An Example

- Specify a different WORKS_ON table

```
CREATE VIEW WORKS_ON_BYNAME AS
SELECT FNAME, LNAME, PNAME, HOURS
      FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE SSN=ESSN AND PNO=PNUMBER
      GROUP BY PNAME;
```

Using a Virtual Table

- We can specify SQL queries on a newly created table (view):

```
SELECT FNAME, LNAME  
FROM WORKS_ON_BYNAME  
WHERE PNAME='Seena' ;
```

- When no longer needed, a view can be dropped:

```
DROP WORKS_ON_NEW ;
```

Efficient View Implementation

- Query modification:
 - Represent the view query in terms of a query on the underlying base tables
- Disadvantage:
 - Inefficient for views defined via complex queries
 - Especially if additional queries are to be applied to the view within a short time period

Efficient View Implementation

- View materialization:
 - Involves physically creating and keeping a temporary table
- Assumption:
 - Other queries on the view will follow
- Concerns:
 - Maintaining correspondence between the base table and the view when the base table is updated
- Strategy:
 - Incremental update

Updating Views

- Update on a view of a single table without aggregate operations:
 - Update may map to an update on the underlying base table
- Views involving joins:
 - An update *may* map to an update on the underlying base relations
 - (This is not always possible)

Non-updatable Views

- Views defined using GROUP BY and aggregate functions are not updateable
- Views defined on multiple tables using joins are generally not updateable
- View on tables that have CHECK constraints: the CHECK must also be added to the definition of a view if the view is to be updated
 - To allow check for updatability and to plan for an execution strategy

Constraints

Previous Constraints

CREATE TABLE Teaching (

ProfId INTEGER,

CrsCode CHAR (6),

Semester CHAR (6),

PRIMARY KEY (*CrsCode*, *Semester*),

FOREIGN KEY (*ProfId*) REFERENCES Professor (*Id*)

ON DELETE NO ACTION

ON UPDATE CASCADE,

FOREIGN KEY (*CrsCode*) REFERENCES Course (*CrsCode*)

ON DELETE SET NULL

ON UPDATE CASCADE)

Table Semantic Constraints

- Used for application dependent conditions
- *Example:* limit attribute values

```
CREATE TABLE Transcript (  
    StudId INTEGER,  
    CrsCode CHAR(6),  
    Semester CHAR(6),  
    Grade CHAR(1),  
    CHECK (Grade IN ('A', 'B', 'C', 'D', 'F')),  
    CHECK (StudId > 0 AND StudId < 1000000000) )
```

- Each row in table must satisfy condition

User-Defined Domains

- Possible attribute values can be specified
 - Using a CHECK constraint, or
 - Creating a new domain
- Domain can be used in several declarations
- Domain is a schema element

```
CREATE DOMAIN Grades CHAR (1)
    CHECK (VALUE IN ('A', 'B', 'C', 'D', 'F'))
CREATE TABLE Transcript (
    ...,
    Grade: Grades,
    ... )
```

Table Constraint Example

- Ensure that managers are paid more than their employees.

```
CREATE TABLE Employee (  
    Id INTEGER,  
    Name CHAR(20),  
    Salary INTEGER,  
    MngrSalary INTEGER,  
    CHECK ( MngrSalary > Salary) )
```

Constraints – Problems

- **Problem 1:**

An empty table always satisfies all CHECK constraints (an idiosyncrasy of the SQL standard)

```
CREATE TABLE Employee (  
    Id INTEGER,  
    Name CHAR(20),  
    Salary INTEGER,  
    MngrSalary INTEGER,  
    CHECK ( 0 < (SELECT COUNT (*) FROM Employee)) )
```

- If Employee is empty, there are no rows on which to evaluate the CHECK condition.

Constraints – Problems

- **Problem 2:**

Inter-relational constraints should be symmetric

```
CREATE TABLE Employee (  
  Id INTEGER,  
  Name CHAR(20),  
  Salary INTEGER,  
  MngrSalary INTEGER,  
  CHECK ((SELECT COUNT (*) FROM Manager) <  
          (SELECT COUNT (*) FROM Employee)) )
```

- Why should constraint be in Employee, rather than Manager?
- What if Employee is empty?

Assertions

- Assertions are schema elements
- Symmetrically specifies an inter-relational constraint
- Applies to entire database (not just the individual rows of a single table)
 - hence it works even if Employee is empty

```
CREATE ASSERTION DontFireEveryone  
CHECK (0 < SELECT COUNT (*) FROM Employee)
```

Designing Assertions

- Specify a query that violates the condition; include that inside a `NOT EXISTS` clause
- Query result must be empty
 - if the query result is not empty, the assertion has been violated

```
CREATE ASSERTION KeepEmployeeSalariesDown  
CHECK (NOT EXISTS(  
    SELECT * FROM Employee E  
    WHERE E.Salary > E.MngrSalary))
```

Assertion Example

- The salary of an employee must not be greater than the salary of the manager of the department that the employee works for.

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK (NOT EXISTS
      (SELECT *
       FROM EMPLOYEE E, EMPLOYEE M,
            DEPARTMENT D
       WHERE E.SALARY > M.SALARY AND
            E.DNO=D.NUMBER AND
            D.MGRSSN=M.SSN) )
```

Assertion Example

```
CREATE ASSERTION NoEmptyCourses
CHECK (NOT EXISTS (
    SELECT * FROM Teaching T
    WHERE NOT EXISTS (
        SELECT * FROM Transcript R
        WHERE T.CrsCode = R.CrsCode
        AND T.Semester = R.Semester)
    ))
```

not the double negative logic:

It is **not true** that there are courses taught that **do not** have students.

Triggers

Triggers

- Triggers are active statements that specify responses to specific conditions.
- A trigger is a schema element

```
CREATE TRIGGER CrsChange  
  AFTER UPDATE OF CrsCode, Semester ON Transcript  
  WHEN (Grade IS NOT NULL)  
  ROLLBACK
```

Trigger Overview

- Element of the database schema
- General form:
 - ON *<event>* IF *<condition>* THEN *<action>*
 - *Event*- request to execute database operation
 - *Condition* - predicate evaluated on database state
 - *Action* – execution of procedure that might involve database updates
- Example:

ON updating maximum course enrollment

IF number registered > new max enrollment limit
THEN deregister students using LIFO policy

Possible Trigger Semantics

- **Activation** - Occurrence of the event
- **Consideration** - The point, after activation, when the *condition* is evaluated
 - **Immediate:** evaluate condition as soon as the event occurs
 - **Deferred:** wait to evaluate the condition at the end of the transaction
 - The *condition* may refer to the database state both before and after the triggering event

Possible Trigger Semantics

- **Execution** – point at which the *action* occurs
 - With deferred consideration, execution is also deferred
 - With immediate consideration, execution can occur immediately after consideration or it can be deferred
 - If execution is immediate, execution can occur before, after, or instead of triggering event.
 - Before triggers adapt naturally to maintaining integrity constraints: violation results in rejection of event.

Possible Trigger Semantics

- **Granularity**

- *Row-level granularity*: change of a single row is an event (a single UPDATE statement might result in multiple events)
- *Statement-level granularity*: events are statements (a single UPDATE statement that changes multiple rows is a single event).

Possible Trigger Semantics

- **Multiple Triggers**

- How should multiple triggers activated by a single event be handled?
 - Evaluate one condition at a time and if true immediately execute action or
 - Evaluate all conditions, then execute actions
- The execution of an action can affect the truth of a subsequently evaluated condition so the choice is significant.

Triggers in SQL

- **Events:**
INSERT, DELETE, or UPDATE statements or changes to individual rows caused by these statements
- **Condition:**
Anything that is allowed in a WHERE clause
- **Action:**
An individual SQL statement or a program written in the language of Procedural Stored Modules (PSM)
(which may contain embedded SQL statements)

Triggers in SQL

- **Consideration:** *Immediate*
 - Condition can refer to the state of the affected row or table before *and* after the event occurs
- **Execution:** *Immediate* – can be before or after the execution of the triggering event
 - Action of a before trigger cannot modify the database
- **Granularity:** Both *row-level* and *statement-level*

Trigger Syntax

```
CREATE TRIGGER trigger-name
  { BEFORE | AFTER }
  { INSERT | DELETE | UPDATE [OF column-name-list] }
  ON table-name
  [ REFERENCING { OLD AS old-tuple-name |
                  NEW AS new-tuple-name |
                  OLD TABLE AS old-table-name |
                  NEW TABLE AS new-table-name } ]
  [ FOR EACH { ROW | STATEMENT } ]
  [ WHEN (precondition) ]
  statement-list
```

Before Trigger (row granularity)

```
CREATE TRIGGER Max_EnrollCheck
  BEFORE INSERT ON Transcript
    REFERENCING NEW AS N  --row to be added
  FOR EACH ROW
  WHEN
    ((SELECT COUNT (T.StudId) FROM Transcript T
      WHERE T.CrsCode = N.CrsCode
        AND T.Semester = N.Semester)
    >=
    (SELECT C.MaxEnroll FROM Course C
      WHERE C.CrsCode = N.CrsCode ))
  ABORT TRANSACTION
```

Check that
enrollment \leq limit

After Trigger (row granularity)

```
CREATE TRIGGER LimitSalaryRaise
AFTER UPDATE OF Salary ON Employee
REFERENCING OLD AS O
              NEW AS N
FOR EACH ROW
WHEN (N.Salary - O.Salary > 0.05 * O.Salary)
  UPDATE Employee
  SET Salary = 1.05 * O.Salary
  WHERE Id = O.Id
```

*No salary raises
greater than 5%*

Note: The action itself is a triggering event (but in this case a chain reaction is not possible)

After Trigger (stmt granularity)

```
CREATE TRIGGER RecordNewAverage  
AFTER UPDATE OF Salary ON Employee  
FOR EACH STATEMENT  
  INSERT INTO Log  
  VALUES (CURRENT_DATE,  
           SELECT AVG (Salary)  
           FROM Employee)
```

*Keep track of salary
averages in the log*

Trigger Example

- A trigger to compare an employee's salary to his/her supervisor during insert or update operations:

```
CREATE TRIGGER INFORM_SUPERVISOR
BEFORE INSERT OR UPDATE OF
    SALARY, SUPERVISOR_SSN ON EMPLOYEE
FOR EACH ROW
WHEN
    (NEW.SALARY > (SELECT SALARY FROM EMPLOYEE
                    WHERE SSN=NEW.SUPERVISOR_SSN))
INFORM_SUPERVISOR(NEW.SUPERVISOR_SSN,NEW.SSN) ;
```

INFORM_SUPERVISOR is a *stored procedure*.