


Advanced Database Management Systems

Lecture 15

Index Structures: Sections 14.1, 14.2

Review: Disks

Review: Hard Drives




What are the primary factors in computing the time for data transfer?

Disk Parameters

parameter		typical value	source
s	seek time	3 msec	fixed
p	rotational velocity	10,000 rpm 167 rps	fixed
rd	rotational delay (latency)	2 msec	$.5 \cdot (1/p)$ (average)
T	track size	50 Kbytes	fixed
B	block size	512-4096 bytes	formatted
G	interblock gap size	128 bytes	formatted
tr	transfer rate	800Kbytes/sec	$T \cdot p$
btt	block transfer time	1 msec	B/tr
btr	bulk transfer rate (consecutive blocks)	700Kbytes/sec	$(B/(B+G)) \cdot tr$

Random Block Transfer Time



rbtt = time to locate and transfer one random block
= $s + rd + btt$ ($3 + 2 + 0.1 = 5.1 \text{ msec}$)


Transferring Multiple Blocks

Time to transfer n blocks of data:



randomly located blocks:

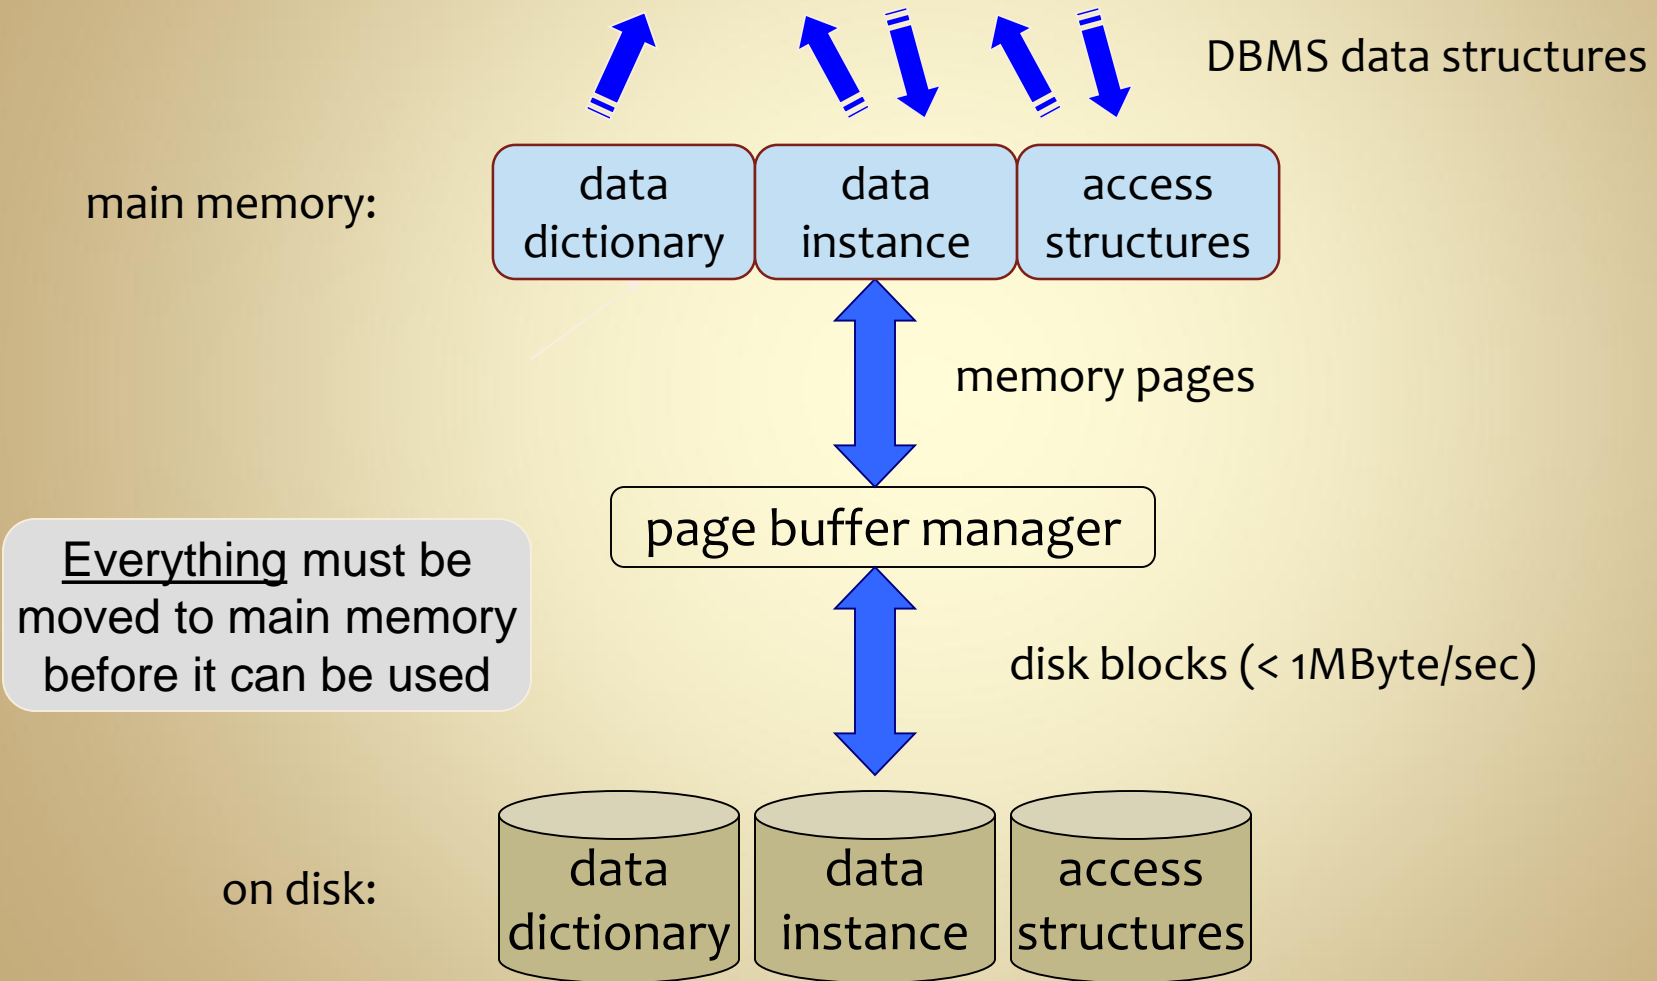
$$rbtt = n \cdot (s + rd + btt)$$



consecutively located blocks:

$$cbtt = s + rd + n \cdot btt$$

Page Buffer Management



Fundamental Results

- **Organize data in blocks**
 - a block is the basic unit of transfer
- **Layout data to maximize possibility of consecutive block retrieval**
 - avoid seek time and latency
- **This will impact**
 - record layout in files
 - access structure (indexes, trees) organization

Review: Files

Blocks, Records and Files

- A **record** is one unit of structured data
 - **tuple** in a relation
 - **node** in a tree or index
 - object or other structure
- Records may be *fixed length* or *variable length*
- A **file** is a set of records stored as a unit on disk
 - a file is stored on some set of disk blocks

File Organization Options

- Contiguous vs. non-contiguous
 - Is the file stored in contiguous blocks, or scattered around the disk
- Ordered vs. non-ordered
 - Are the records ordered on some ordering field?
- Spanning vs. non-spanning
 - Can a record be split across blocks, or is every record contained within one block?

Blocking Factors

- *Blocking factor* determines how many records can fit in a block
 - $bfr = \lfloor B / R \rfloor$ where B = block size and R = record size
- *Number of blocks* required is determined by the blocking factor and the number of records
 - $b = \lceil r / bfr \rceil$ where r = number of records

File Structure

- Some navigational structure is needed to identify/locate all blocks that comprise a file
 - Similar to dynamic memory management
- Possible structures:
 - **contiguous**: Once you find the first block, keep reading blocks in sequence until the end
 - **linked**: Each block has a pointer to the next block
 - **indexed**: an access structure (index or tree) holds block pointers to all blocks in the file

Unordered vs. Ordered Files

- Unordered

- insert: new records placed at the end of file
 - $O(1)$
- find/delete: linear search
 - $O(n)$

- Ordered

- insert: requires reorganization
 - $O(n)$
- find/delete on ordering field: binary search
 - $O(\log n)$
- find/delete on other field: linear search
 - $O(n)$

Ordered File Example

	NAME	SSN	BIRTHDATE	JOB	SALARY	SEX
block 1	Aaron, Ed					
	Abbott, Diane					
	⋮					
	Acosta, Marc					
block 2	Adams, John					
	Adams, Robin					
	⋮					
	Akers, Jan					
block 3	Alexander, Ed					
	Alfred, Bob					
	⋮					
	Allen, Sam					
block 4	Allen, Troy					
	Anders, Keith					
	⋮					
	Anderson, Rob					
block 5	Anderson, Zach					
	Angeli, Joe					
	⋮					
	Archer, Sue					


block 6	Arnold, Mack					
	Arnold, Steven					
	⋮					
	Atkins, Timothy					
⋮						
block n-1	Wong, James					
	Wood, Donald					
	⋮					
	Woods, Manny					
block n	Wright, Pam					
	Wyatt, Charles					
	⋮					
	Zimmer, Byron					

Disk Access Costs

rbtt = random block transfer time = $s + rd + btt$ (bad)

cbtt = consecutive block transfer time = btt (good)

	insert	find/delete
unordered non-contiguous	$O(1) \cdot rbtt$	$O(n) \cdot rbtt$
unordered contiguous	$O(1) \cdot rbtt$	$O(n) \cdot cbtt$
ordered non-contiguous	$O(n) \cdot cbtt$	$O(\log n) \cdot rbtt$



assumes we don't reorganize file,
simply mark the record as deleted

Access Structures

Access Structures

- Access structures are auxiliary data structures constructed to assist in locating records in a file
- Benefit of access structures
 - more efficient data access
- Cost of access structures
 - storage space:
access structures must persist along with the data
 - update time:
access structures must be modified when file is modified

Index Blocking

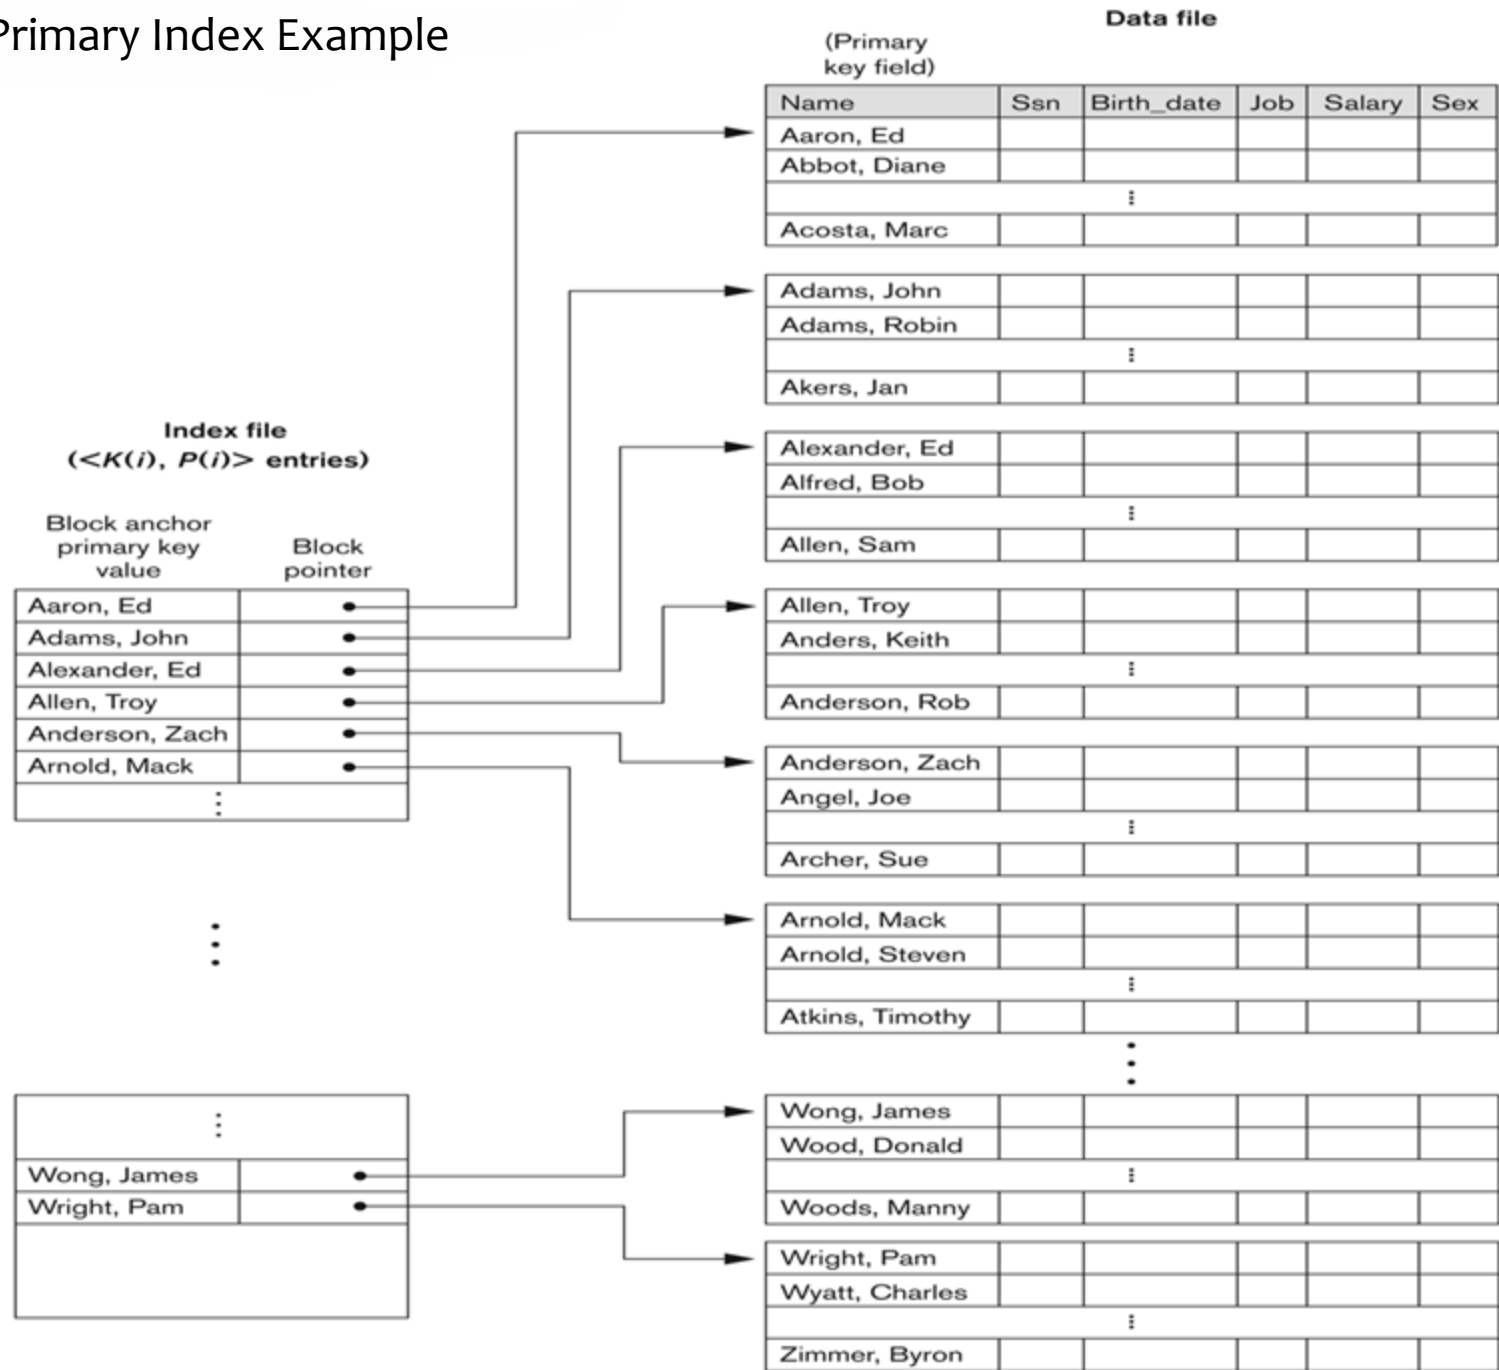
- Access structures also have to be blocked
 - records are index nodes
- *Blocking factor*: nodes per block
 - $bfr_i = \lfloor B/R_i \rfloor$ where B = block size and R_i = node size
- Blocks required to store index
 - $b_i = \lceil r_i / bfr_i \rceil$ where r_i = number of nodes

Single-Level Indexes

Primary Index

- Primary index:
 - fixed length records, non-spanning
 - each record has two fields: < key, block pointer >
 - index file is ordered on key
 - one index record for each block in data file
 - key must be the ordering field of the data file
 - only hold key values for first record (anchor) of each data file block

Primary Index Example



Primary Indexes

- Advantage of a Primary Index
 - much smaller than data file
 - index records (usually) smaller than data records
 - number of index records smaller than number of data records
- Lookup algorithm:
 - binary search on index, then read data block
 - cost with index: $\log_2(r_i) + 1$, $r_i = \# \text{ index blocks}$
 - cost without index: $\log_2(r)$, $r = \# \text{ data blocks}$

Primary Index Example

Data File

data records $r = 200,000$ records

record size $R = 90$ bytes

block size $B = 1024$ bytes

key size $V = 9$ bytes

block pointer size $P = 6$ bytes

data file blocking factor $bfr = \lfloor B/R \rfloor = 11$ records/block

data file blocks $b = \lceil r / bfr \rceil = 18,182$ blocks

binary search cost without index (worse case):

$$\log_2(b) = \log_2(18,182) = 15$$

Primary Index Example

Index File

index records $r_i = b = 18,182$ records

block size $B = 1024$ bytes

key size $V = 9$ bytes

block pointer size $P = 6$ bytes

index record size $R_i = V + P = 15$ bytes

index file blocking factor $bfr_i = \lfloor B / R_i \rfloor = 68$ records / block

index file blocks $b_i = \lceil r_i / bfr_i \rceil = 268$ blocks

binary search cost with index (worse case):

$$\log_2(b) + 1 = \log_2(268) + 1 = 10$$

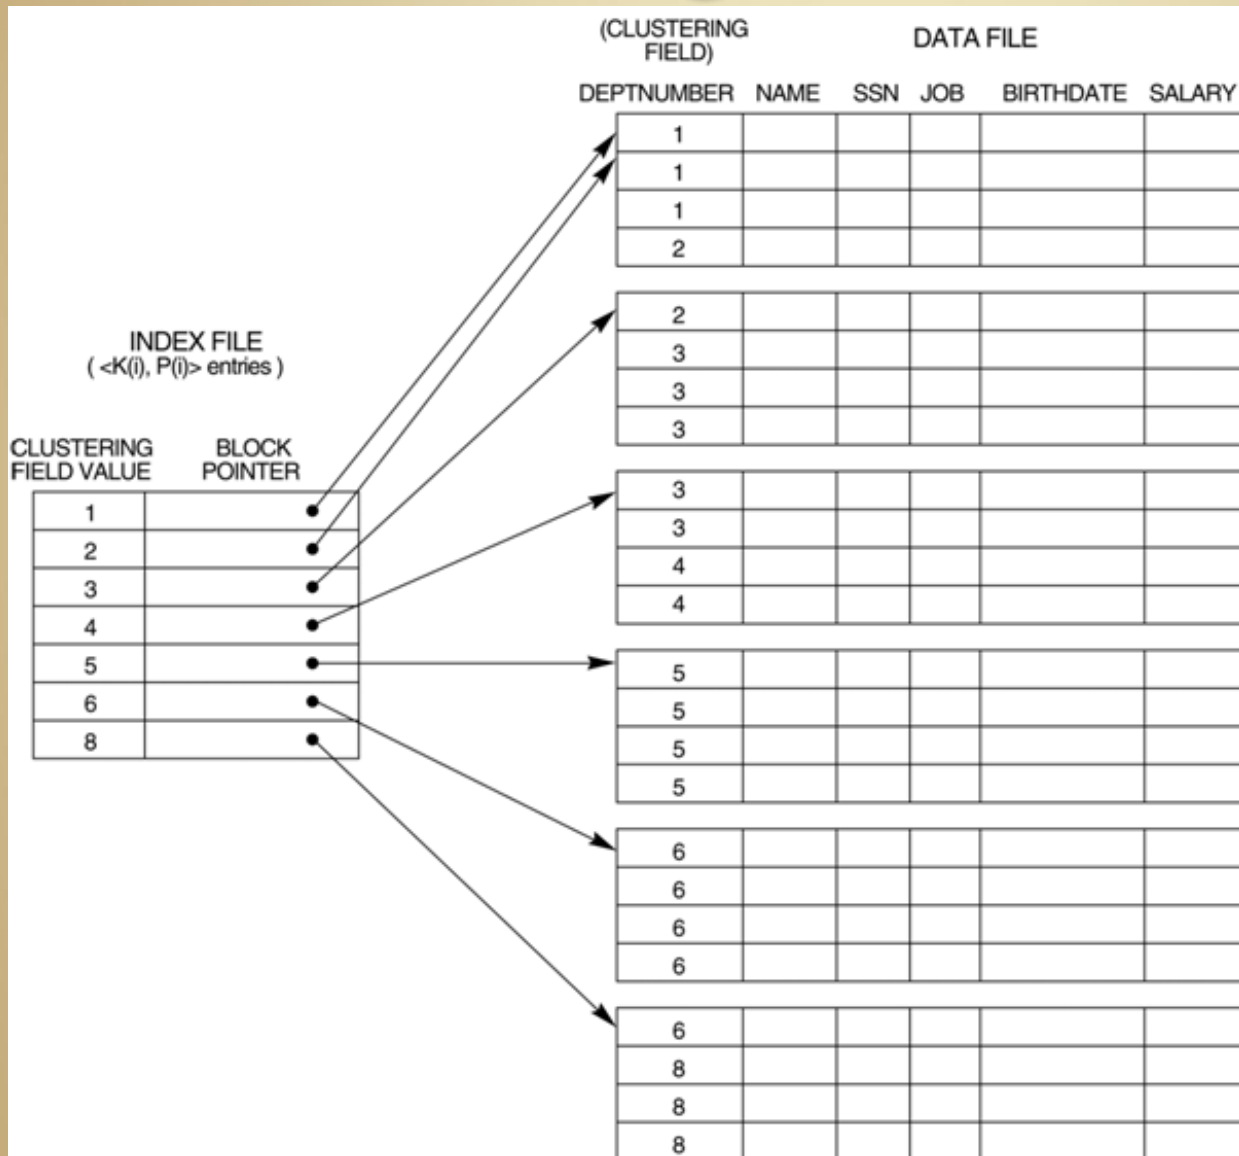
Index Classifications

- **dense index:**
index entry for every data record
- **sparse index:**
index entries for subset of data records
(block anchors)
- **primary index:** key, ordering field
 - sparse
- **clustering index:** non-key, ordering field
 - sparse
- **secondary index:** non-ordering field
 - dense or sparse

Clustering Index

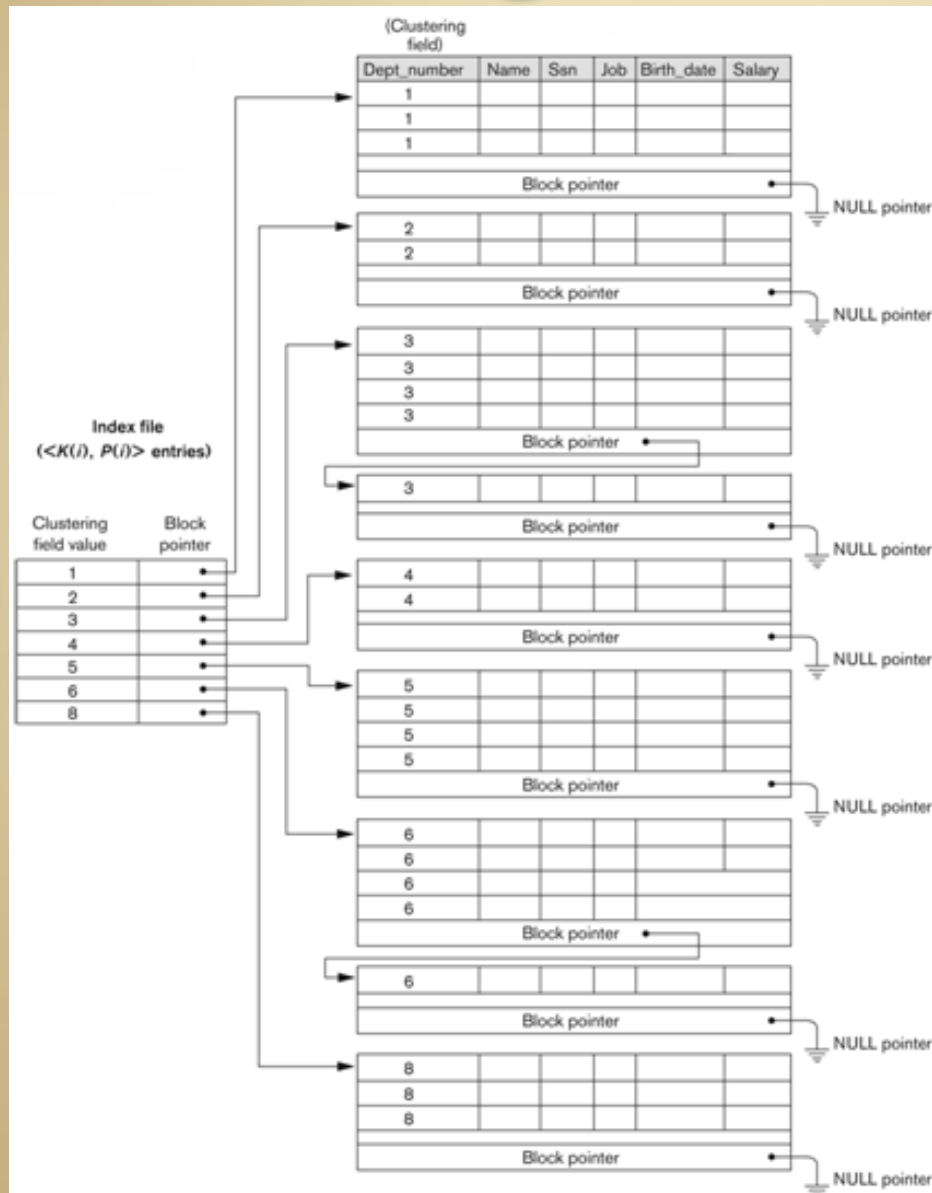
- Data file is ordered on a non-key (non-unique) field
- One index entry for *each distinct value* of the field
 - index entry points to the first data block that contains records with that field value.

Clustering Index Example



assumes
contiguous storage
of data file

Clustering Index Example



non-contiguous storage
of data file

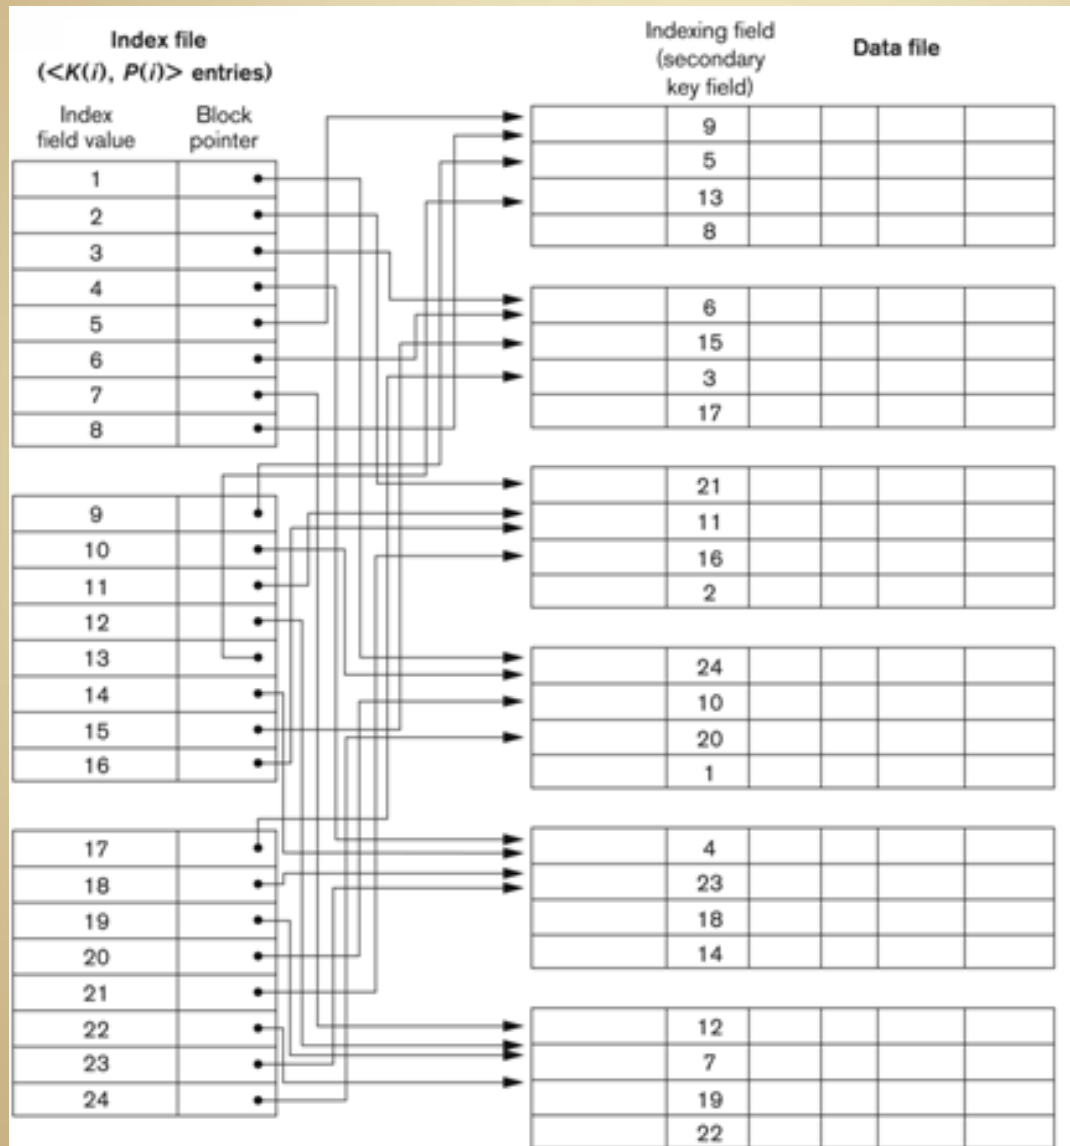
blocks are linked by
block pointers

all records in same data
block have same
clustering field value

Secondary Indexes

- Index on a non-ordering field
 - must be a dense index: one entry for every data record
 - indexing field may be key or non-key
- There may be multiple secondary indexes for the same data file
- Advantage of a secondary index:
 - provides a structure in which index field is ordered, thus allowing binary search

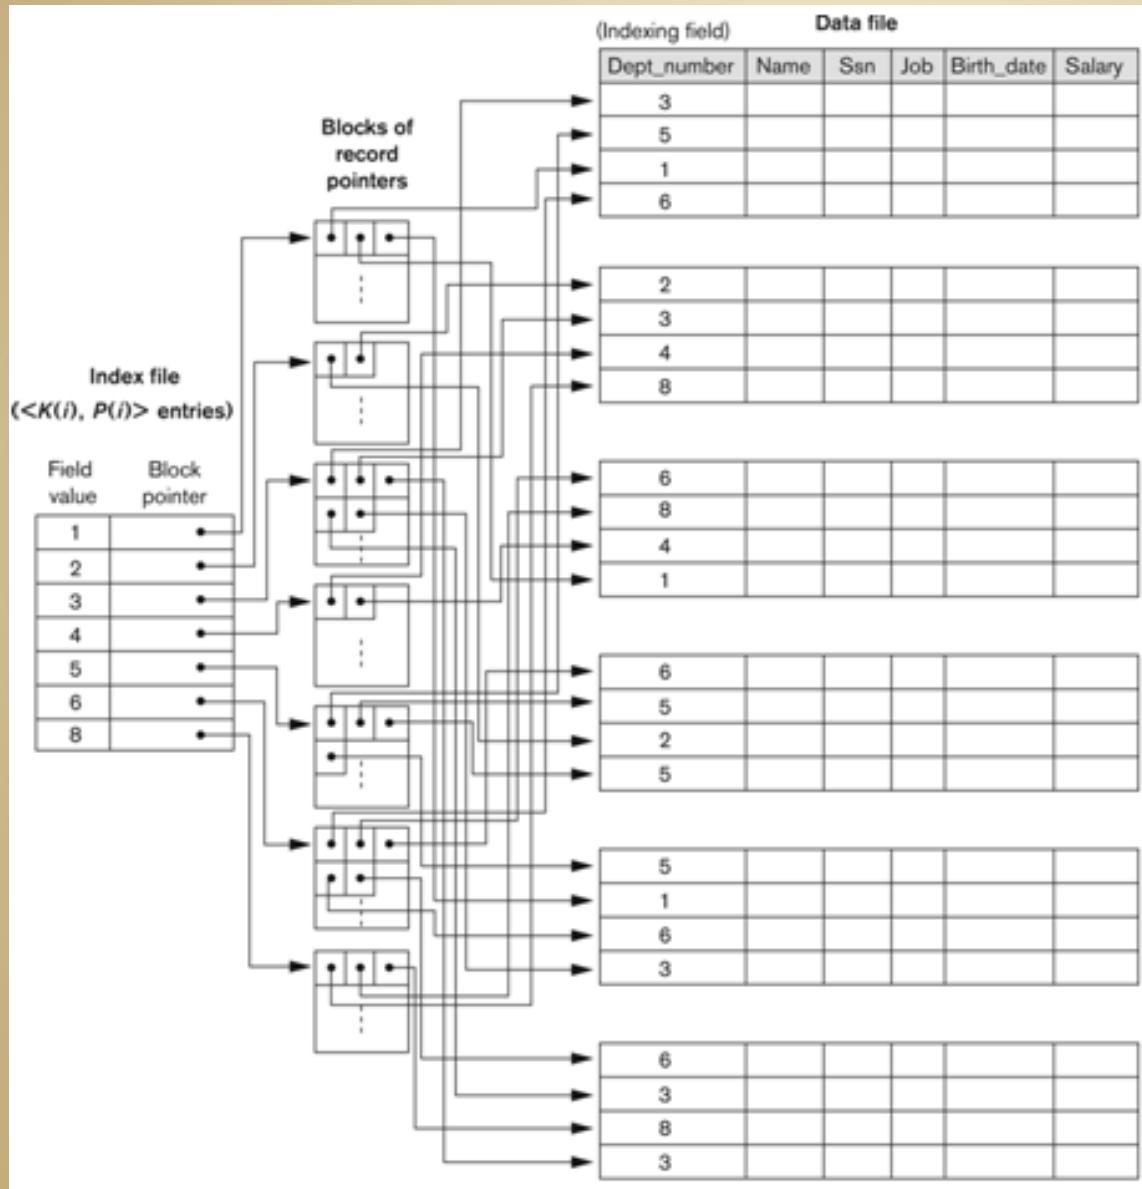
Secondary Index Example



key,
non-ordering
index field

binary search
on index gives
data block of
target record

Secondary Index Example



non-key,
non-ordering
index field

may not be effective
when a small
number of index
values are evenly
distributed among
data blocks

Properties of Index Types

TABLE 14.2 PROPERTIES OF INDEX TYPES

TYPE OF INDEX	NUMBER OF (FIRST-LEVEL) INDEX ENTRIES	DENSE OR NONDENSE	BLOCK ANCHORING ON THE DATA FILE
Primary	Number of blocks in data file	Nondense	Yes
Clustering	Number of distinct index field values	Nondense	Yes/no ^a
Secondary (key)	Number of records in data file	Dense	No
Secondary (nonkey)	Number of records ^b or Number of distinct index field values ^c	Dense or Nondense	No

^aYes if every distinct value of the ordering field starts a new block; no otherwise.

^bFor option 1.


^cFor options 2 and 3.

Multi-level Indexes

Multi-level Primary Indexes

- multi-level index:
a tree built by indexing the indexes
- tree arity (fan-out) is the index blocking factor
- tree height is $\log_{fo}(b)$
- Example: $b = 20,000$ data blocks, $bfr_i = fo = 60$
 - first level index: $b_1 = \lceil 20000/60 \rceil = 334$ blocks
 - second level index: $b_2 = \lceil 334/60 \rceil = 6$ blocks
 - third level index: $b_3 = \lceil 6/60 \rceil = 1$ block
 - lookup cost:
one index block read per level + one data block read
= 4 (random) block reads


Multi-level Primary Index



Multilevel Secondary Indexes

- Since the first level index is ordered, it can be indexed in the same manner as primary key indexes

Multi-level Secondary Index



Static vs. Dynamic Indexes

- Indexes seen so far are static
 - built from a fixed data file
- Updates to data file require rebuilding the index
 - could be very expensive
- Solution: Dynamic Indexes
 - B-trees and B⁺-trees can be adjusted as the data file changes