

Advanced Database Management Systems

Lecture 16
Dynamic Indexes : Sections 14.3

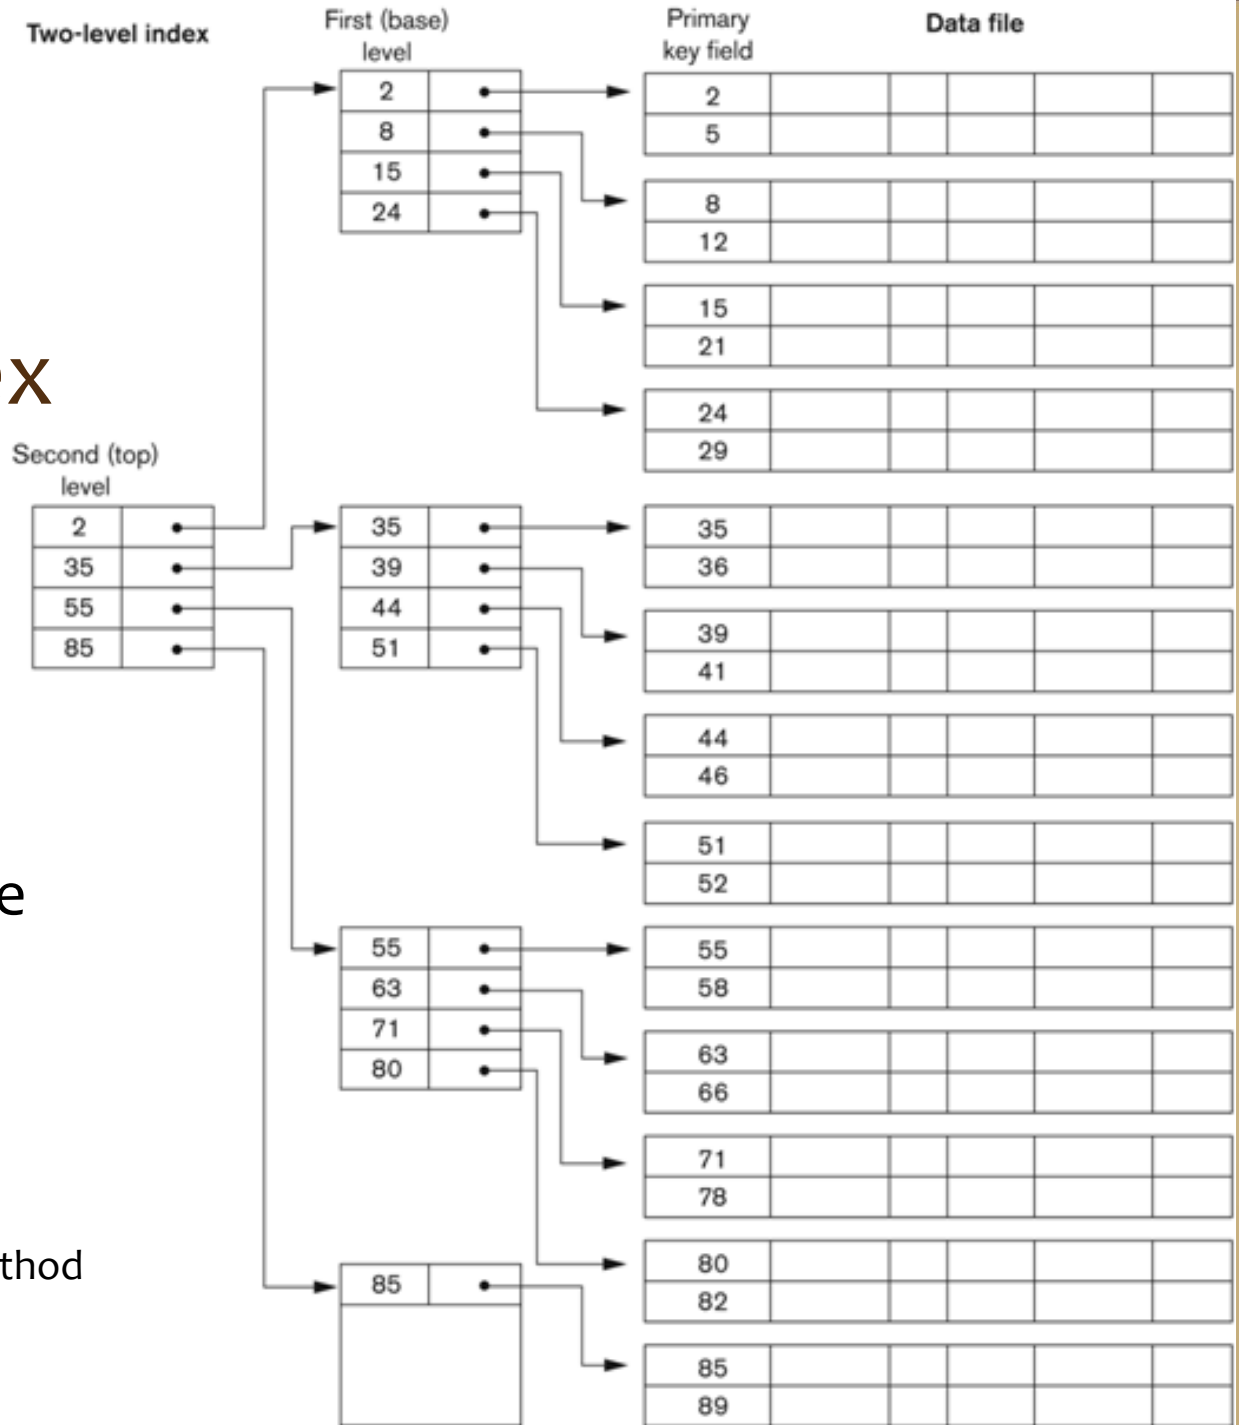
Multi-Level Indexes

- Since a single-level index is an ordered file, we can create a primary index *to the index itself*
 - In this case, the original index file is called the *first-level index* and the index to the index is called the *second-level index*.
- We can repeat the process
 - Create additional levels until all entries of the *top level* fit in one disk block
- A multi-level index can be created for any type of first-level index (primary, secondary, clustering)

Two-level, static, primary index

This is similar to the
ISAM organization
used in early IBM
systems

Index Sequential Access Method



Schematic view of multilevel index


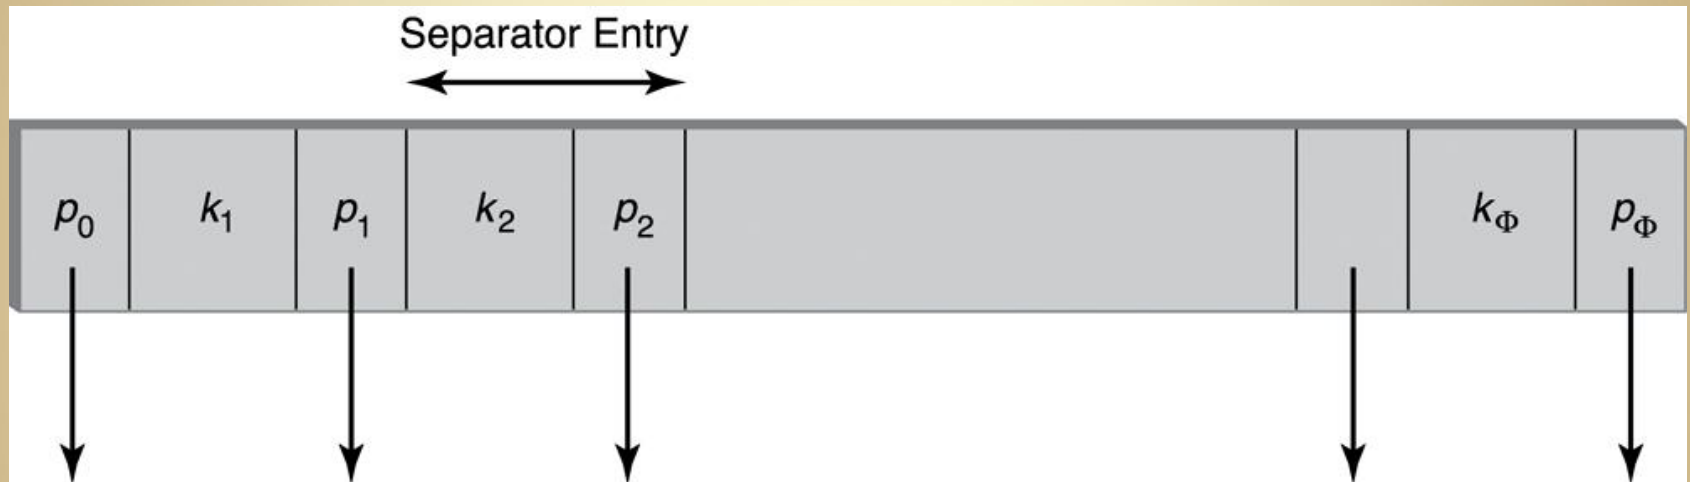
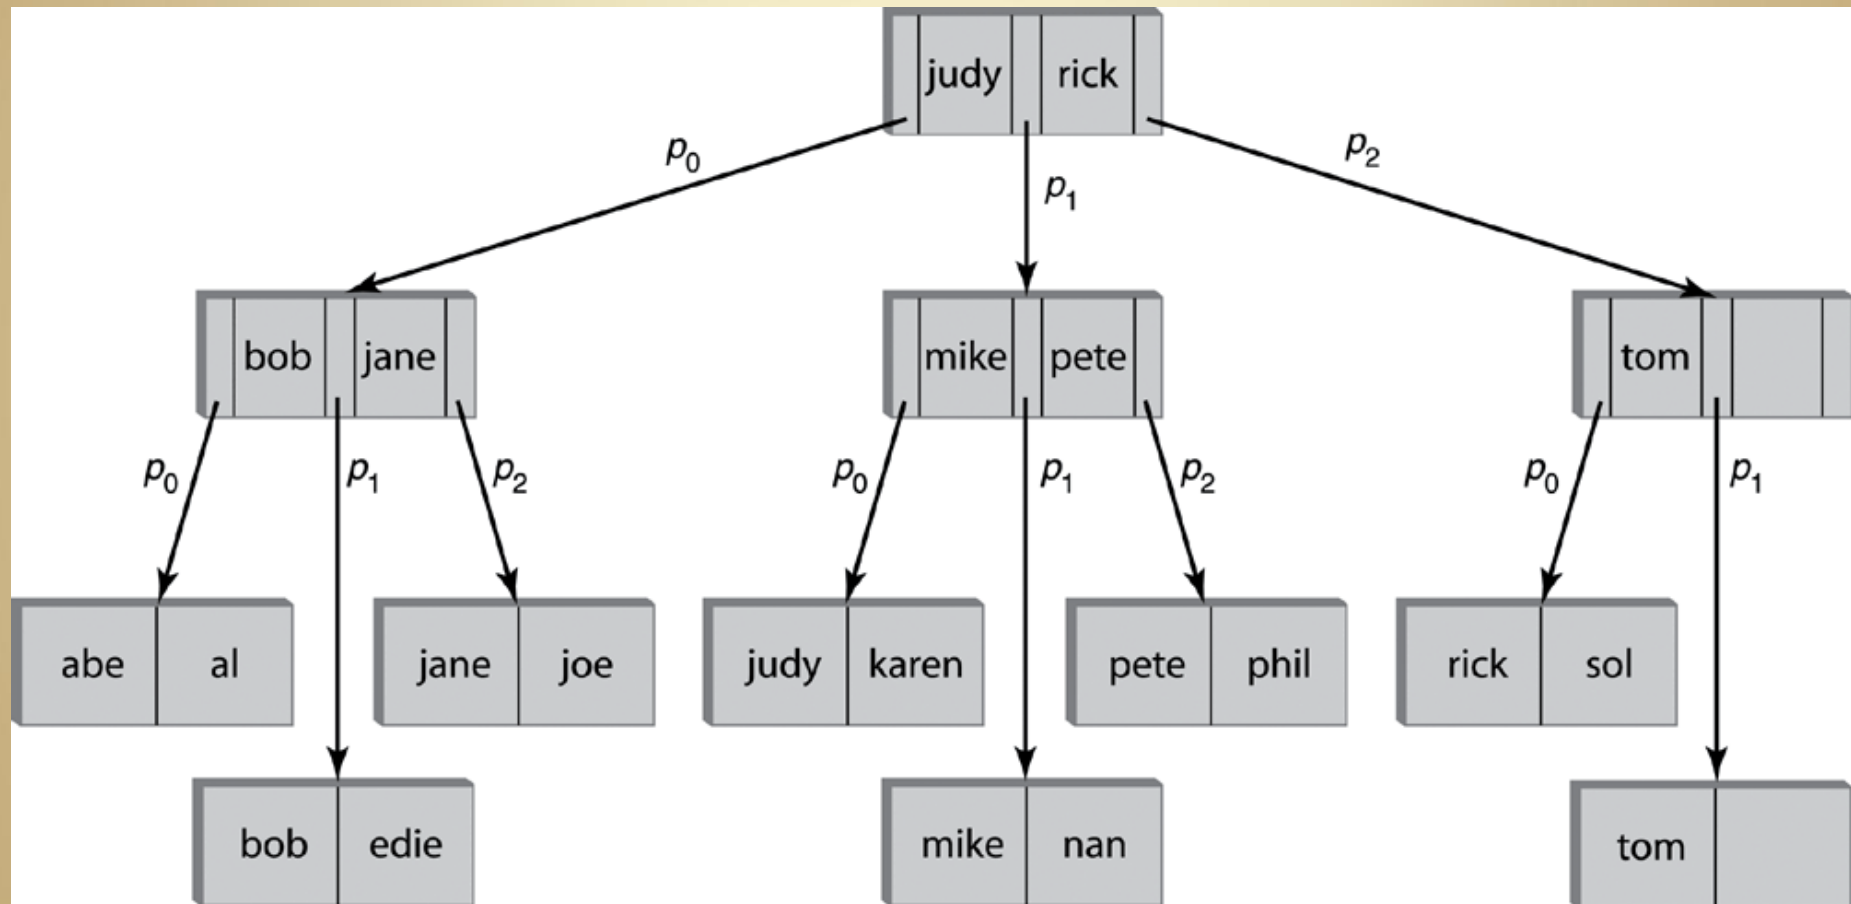


Figure 9.15 Page at a separator level in an ISAM index.



Example of an ISAM index



Dynamic Indexes

- Previous indexes are *static*
 - built up level by level from a data instance (always balanced)
 - must be rebuilt if the record set changes
- Multi-level indexes are a form of *search tree*
 - for static indexes, insertion and deletion of new index entries is a severe problem, since every level of the index is an *ordered file*.

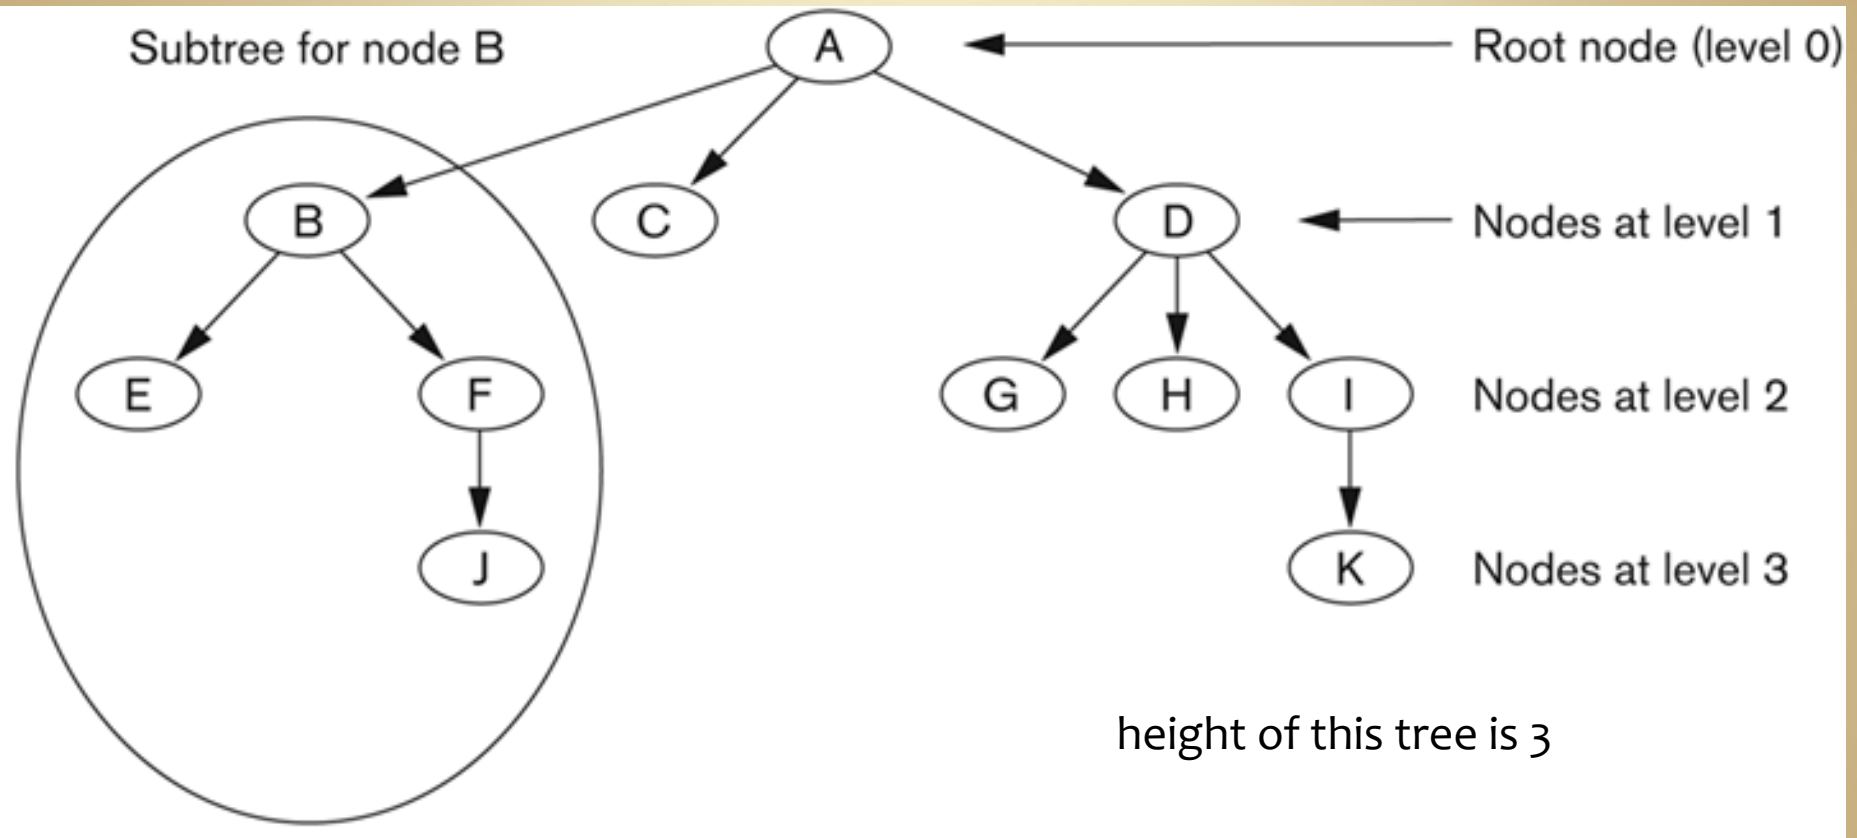
Dynamic Indexes

- *Dynamic indexes* are modified as the record set changes
- *Dynamic indexes* are *balanced search tree* structures
 - tree nodes are sized to match block size
 - nodes are not kept in contiguous blocks
 - nodes are left partially filled to avoid extensive modification to the tree when records are inserted
 - insert and delete algorithms are designed to keep tree balanced
 - balanced = all leaf nodes at same level

Trees: Review

- A tree is a rooted, directed, acyclic graph
- every node has 0-1 parents and 0- p children
 - p is the order or fan-out of the tree
- root node has no parent: there is exactly one root
- *leaf* nodes have no children
- *interior* nodes have at least one child
- *height* of a tree is length of longest path
 - (from root to some leaf)
- binary trees (order 2) are typical for in-memory algorithms
- order of disk-based trees is selected to match node size to disk block size


An Unbalanced Tree



balanced tree = path length from root to any leaf node is the same

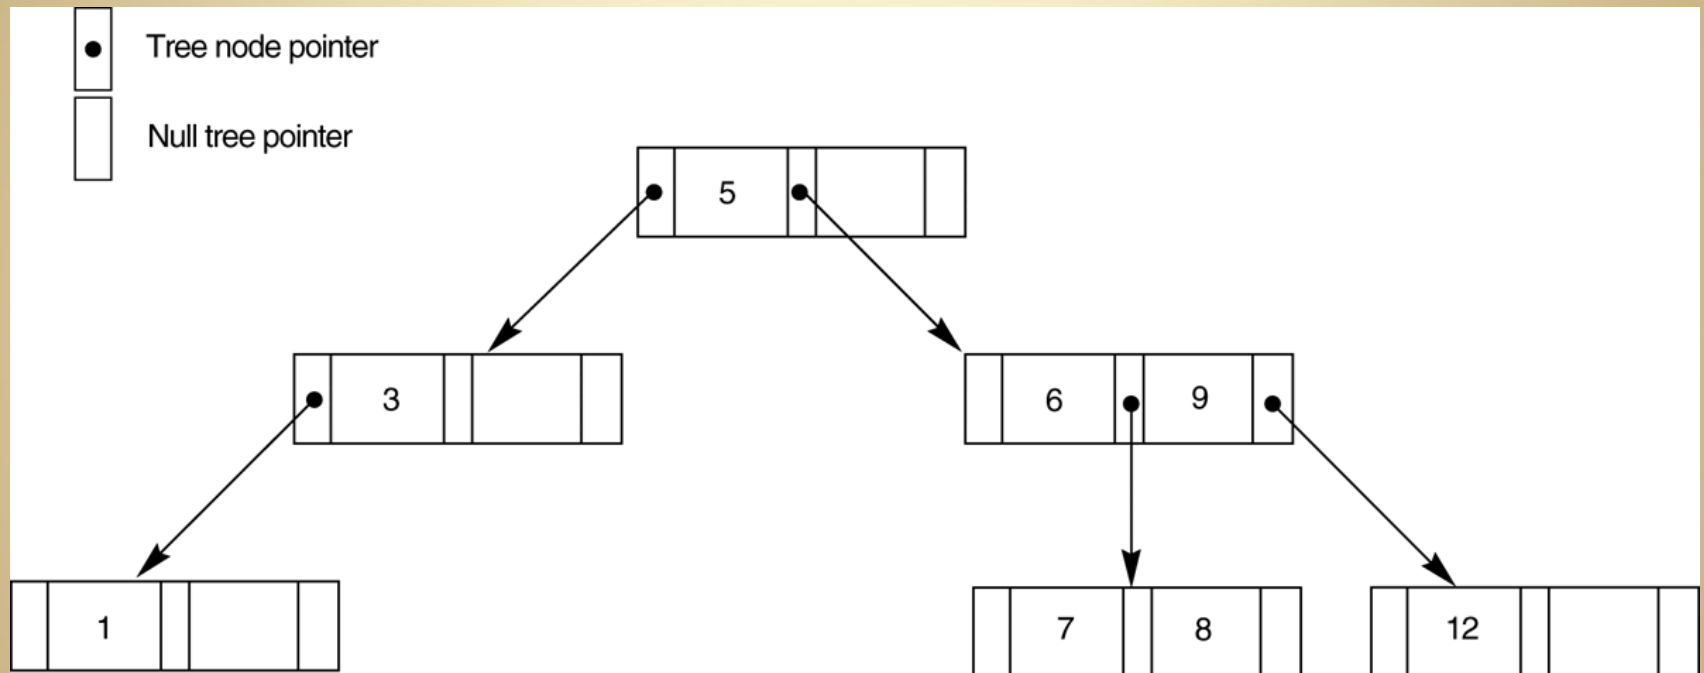
Search Tree: Basic Structure

- Disk based trees have a large fan-out (order)
- Maximize order such that nodes fit in disk blocks/pages
- Following node has order q



A node with order q , will have $q-1$ separator values

Search Tree, Order = 3



Is this a balanced tree?

B-Trees and B+-Trees

- Most multi-level indexes use *B-tree* or *B⁺-tree* data structures
 - Efficient insertion and deletion
 - Each tree node corresponds to a disk block
 - Update algorithms maintain balanced tree
 - Nodes are kept between half-full and completely full to allow for new index entries

Updates to B-Trees & B+-Trees

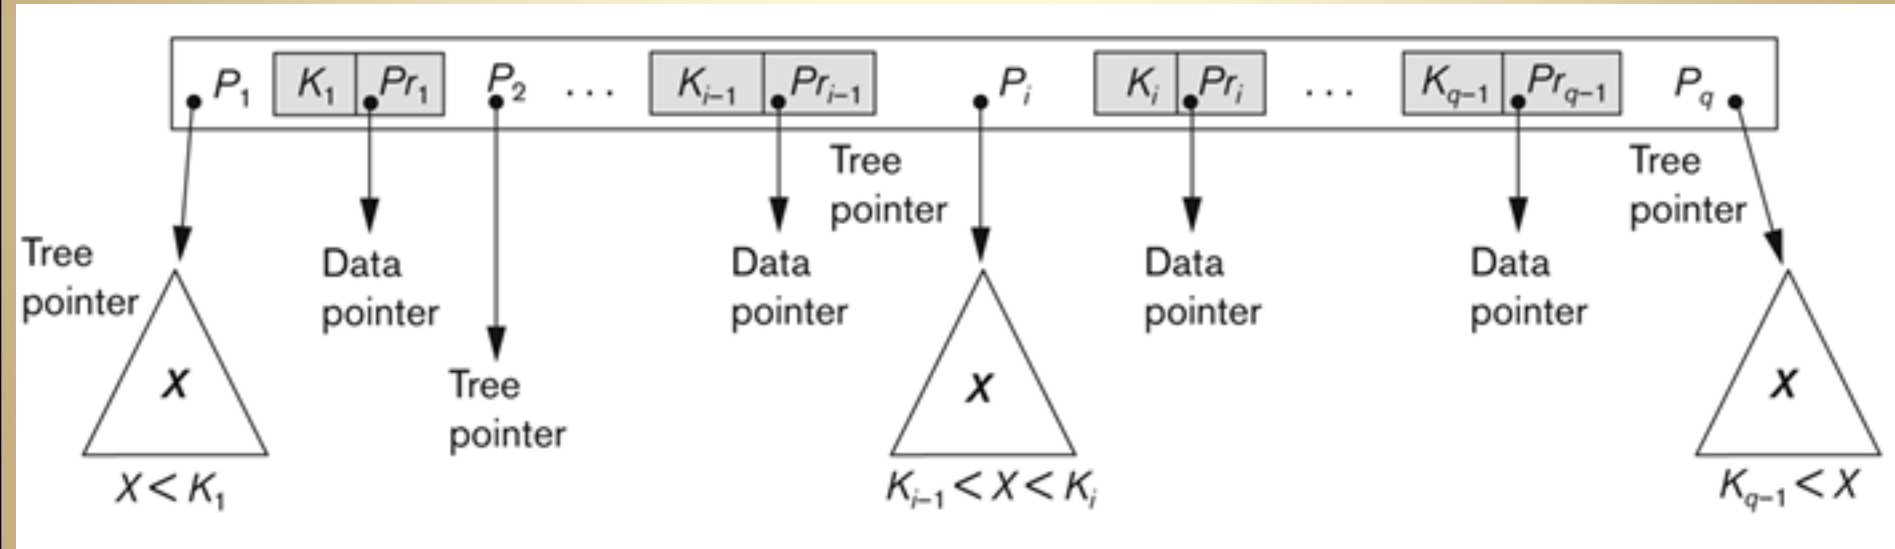
- Insertion into a node that is not full is quite efficient
- Insertion into a full node causes a split into two nodes
 - Splitting may propagate to other tree levels
- Deletion into from a node that is more than half full is quite efficient
- When deletion causes a node to become less than half full, it must be merged with neighboring nodes
 - merging may propagate to other tree levels

B-trees

- B-trees have a single node type
 - interior and leaf nodes have the same structure
 - interior nodes have keys, subtree pointers and data pointers
 - leaf nodes have keys and data pointers, all subtree pointers are NULL
 - **key values appear in exactly one node**


B-tree Structure (order q)

B-tree node with $q-1$ search values and q pointers



Example B-tree

A B-tree of order 3




B-tree, B⁺-tree Comparison

- B-tree:
 - pointers to data records exist at all levels of the tree
 - each index value appears in exactly one node
- B⁺-tree:
 - pointers to data records exists only in the leaf nodes
 - each index value appears in exactly one leaf node
 - some index values also appear in interior nodes

B⁺-Trees

- Support equality and range searches, multi-attribute keys and partial key searches
- Either a secondary index (in a separate file) or the basis for an integrated storage structure
- Responds to dynamic changes in the table
- B⁺-trees have two kinds of nodes
 - *interior* nodes contain keys and *subtree* pointers
 - *leaf* nodes contain keys and *data* pointers
 - key values may appear in multiple nodes


B⁺ Tree Structure



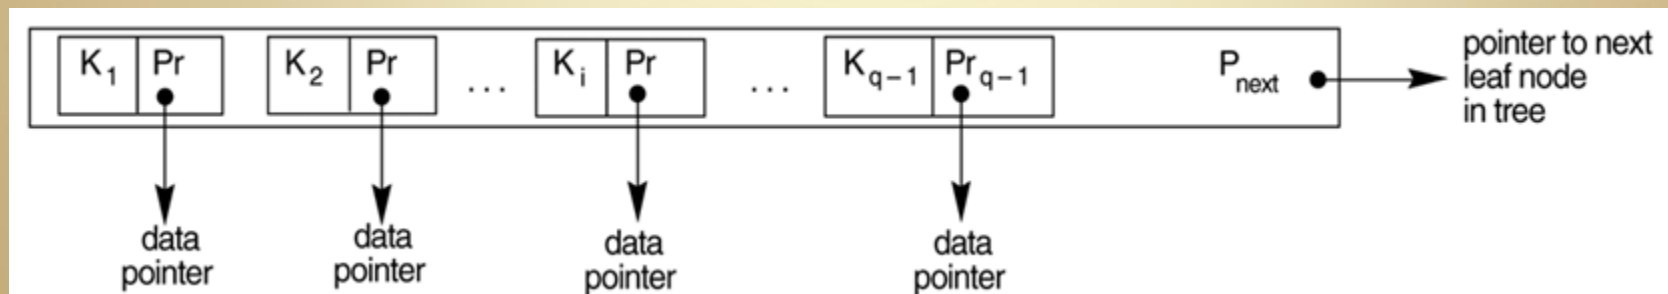
- Leaf level is an ordered linked list of index entries
- All data pointers are in leaves, interior nodes have sub-tree pointers
- Sibling pointers support range searches

B⁺-tree Nodes (Order = q)

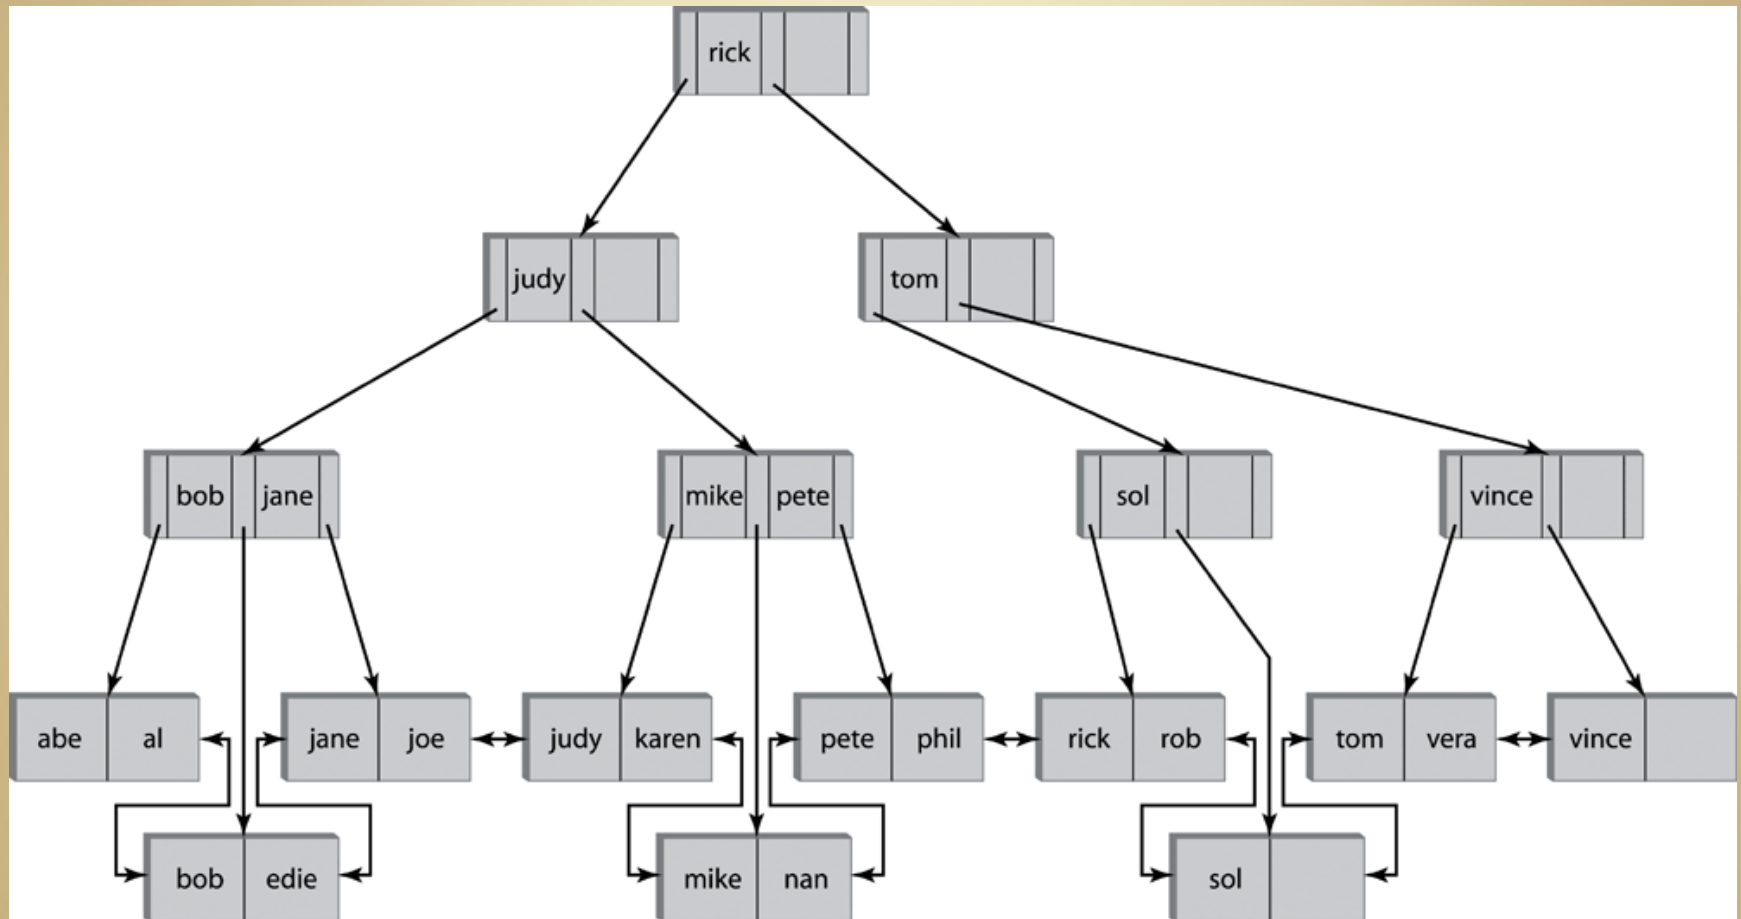
Internal node: $q-1$ search values, q subtree pointers



Leaf node: $q-1$ search values, $q-1$ data pointers,
1 or 2 next node pointers



Example B⁺-tree (order 3)



not shown: each entry in leaf level has a pointer to the data file

Order of a B⁺-tree

- Let B = block size, P = block pointer size and V = index value size.

$$(p - 1)V + pP < B$$

$$p(V + P) - V < B$$

$$p(V + P) < B + V$$

$$p < \frac{B + V}{V + P}$$

$$p = \left\lfloor \frac{B + V}{V + P} \right\rfloor$$

- B⁺-tree nodes are deliberately kept partially filled
 - $q < p$ is the *effective fan-out* of a node
 - q is typically defined as some *fill factor*,
for example: $q = 0.7p$

B⁺-tree Size

- To estimate the size of a B⁺-tree:
 - determine p , from block size, pointer size and key size
 - determine q , from p and desired fill factor
 - determine number of leaf nodes, r_{leaf} , using $q-1$ and number of data records (r)
 - determine number of interior nodes, r_i , at each level using number of blocks in lower level (r_{i-1}) and q
 - When $r_i = 1$, you've reached the root.
The tree height is $h = i$



$$r_0 = r_{leaf} = \left\lceil \frac{r}{q-1} \right\rceil$$

$$r_i = \left\lceil \frac{r_{i-1}}{q} \right\rceil$$

- Blocks required to store the tree is $\sum_{i=0}^h r_i$

B⁺-tree Insertion Example

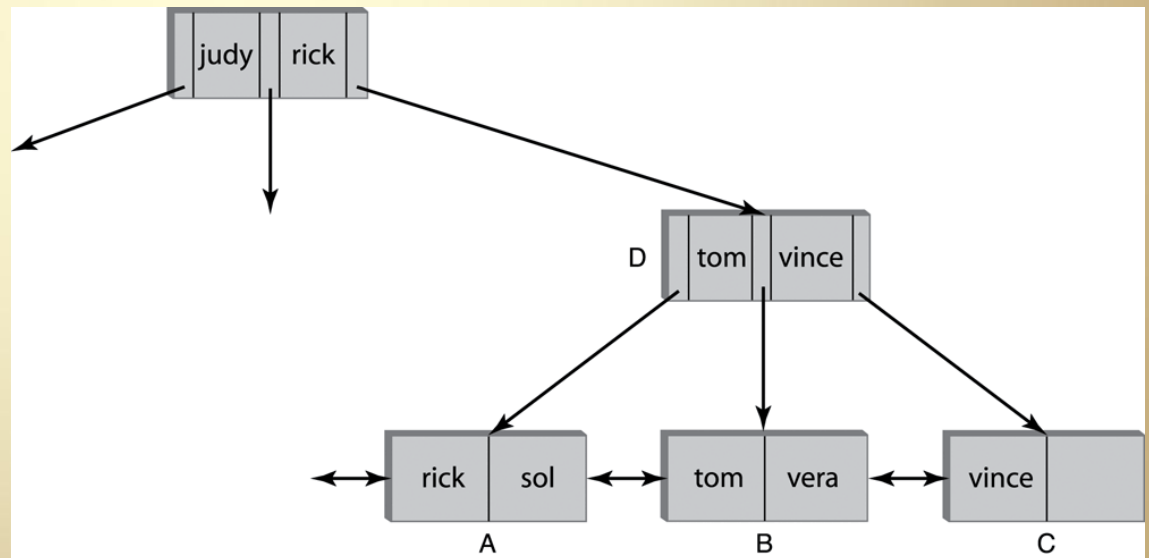

insert "vince"



B⁺-tree Insertion Example

Insert “vera”: Since there is no room in leaf page:

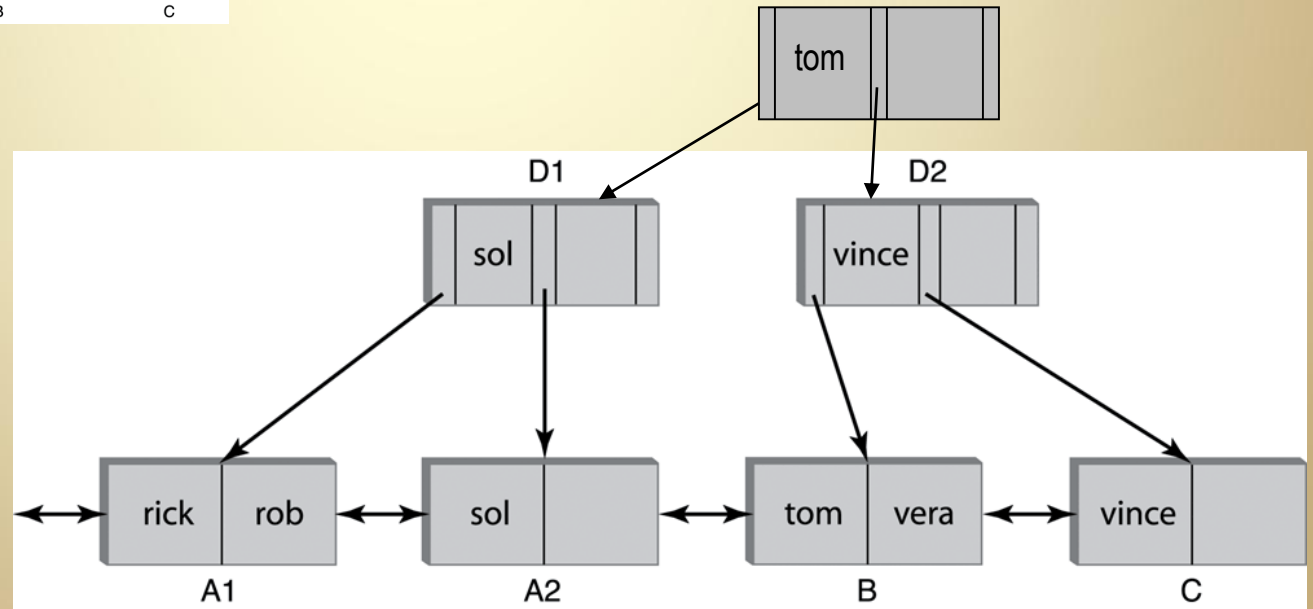
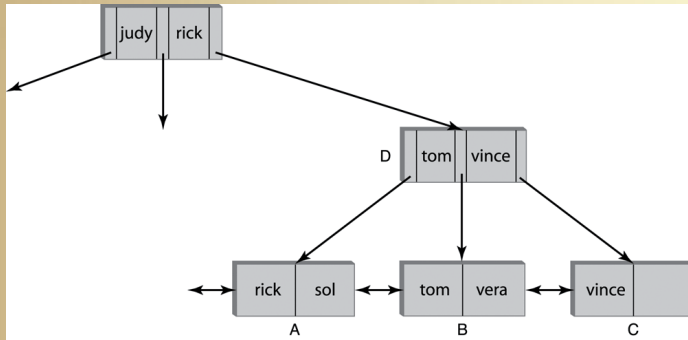
1. Create new leaf, C
2. Distribute index entries between B and C (maintain order)
3. Adjust parent node D



B⁺-tree Insertion Example

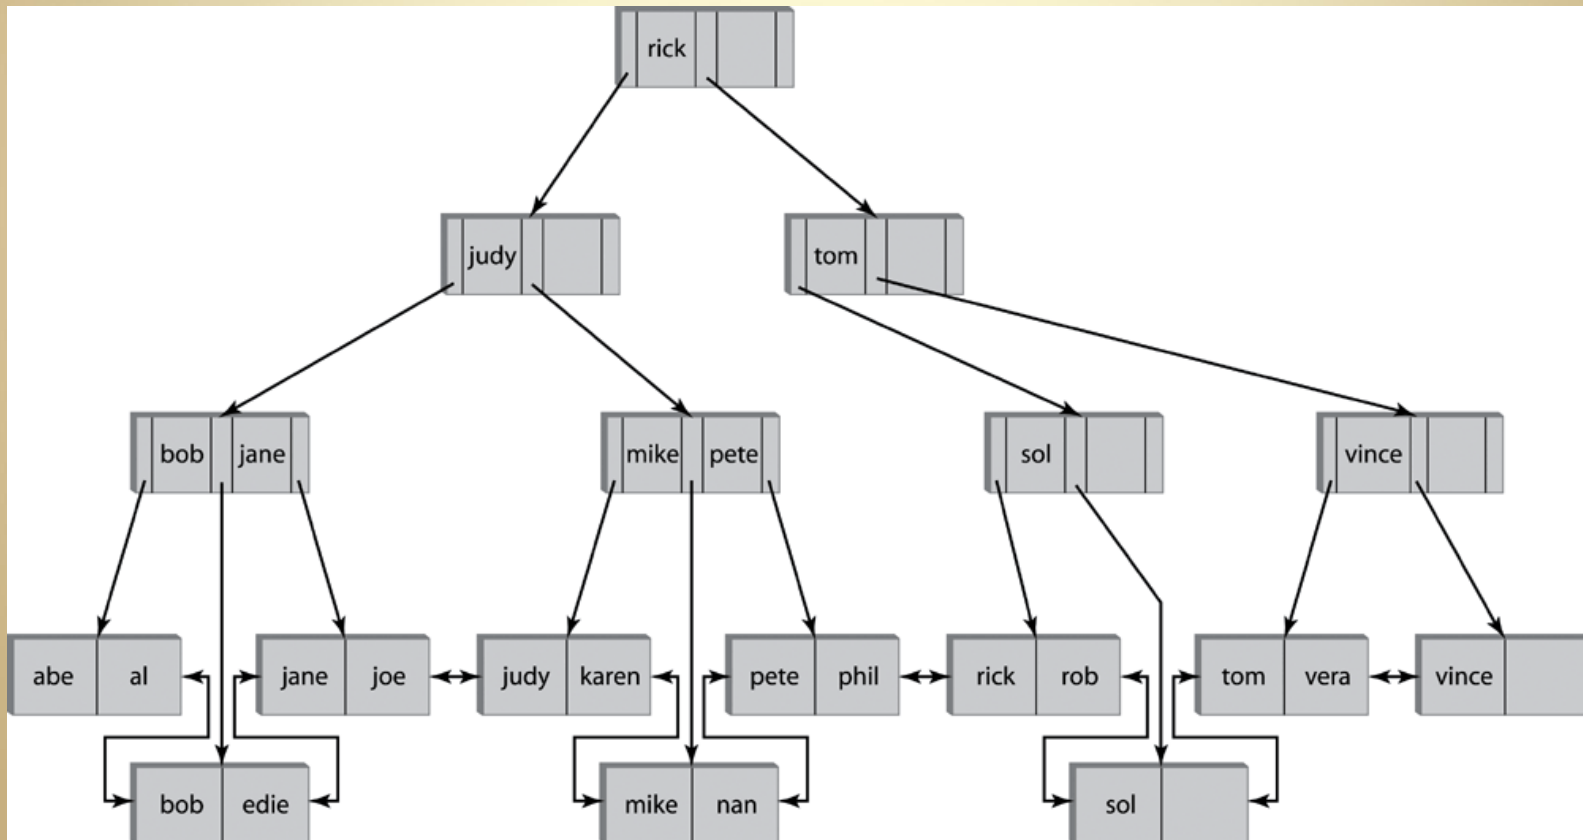
Insert “rob”. Since there is no room in leaf A:

1. Split A into A1 and A2 and distribute index entries between the two (maintain order)
2. Split parent D into D1 and D2 to make room for additional pointer
3. Add new level to connect D1 and D2







B⁺-tree Insertion: Split Propagation



- When splitting an interior node, push a separator up
- Repeat process at next level
- Height of tree increases by one





Example: B⁺-tree Insertion



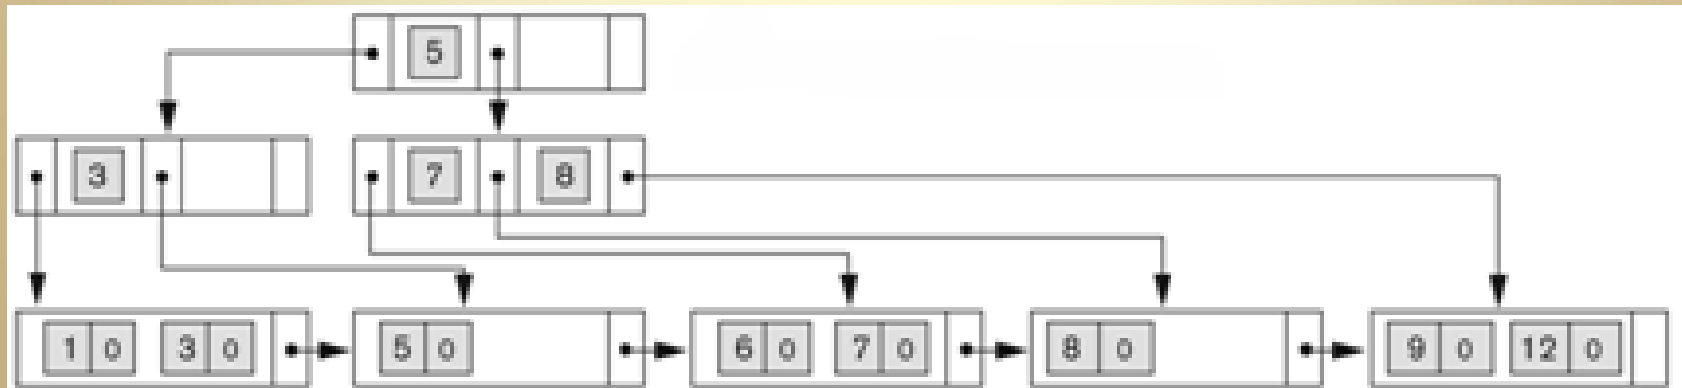

Example: B⁺-tree Insertion



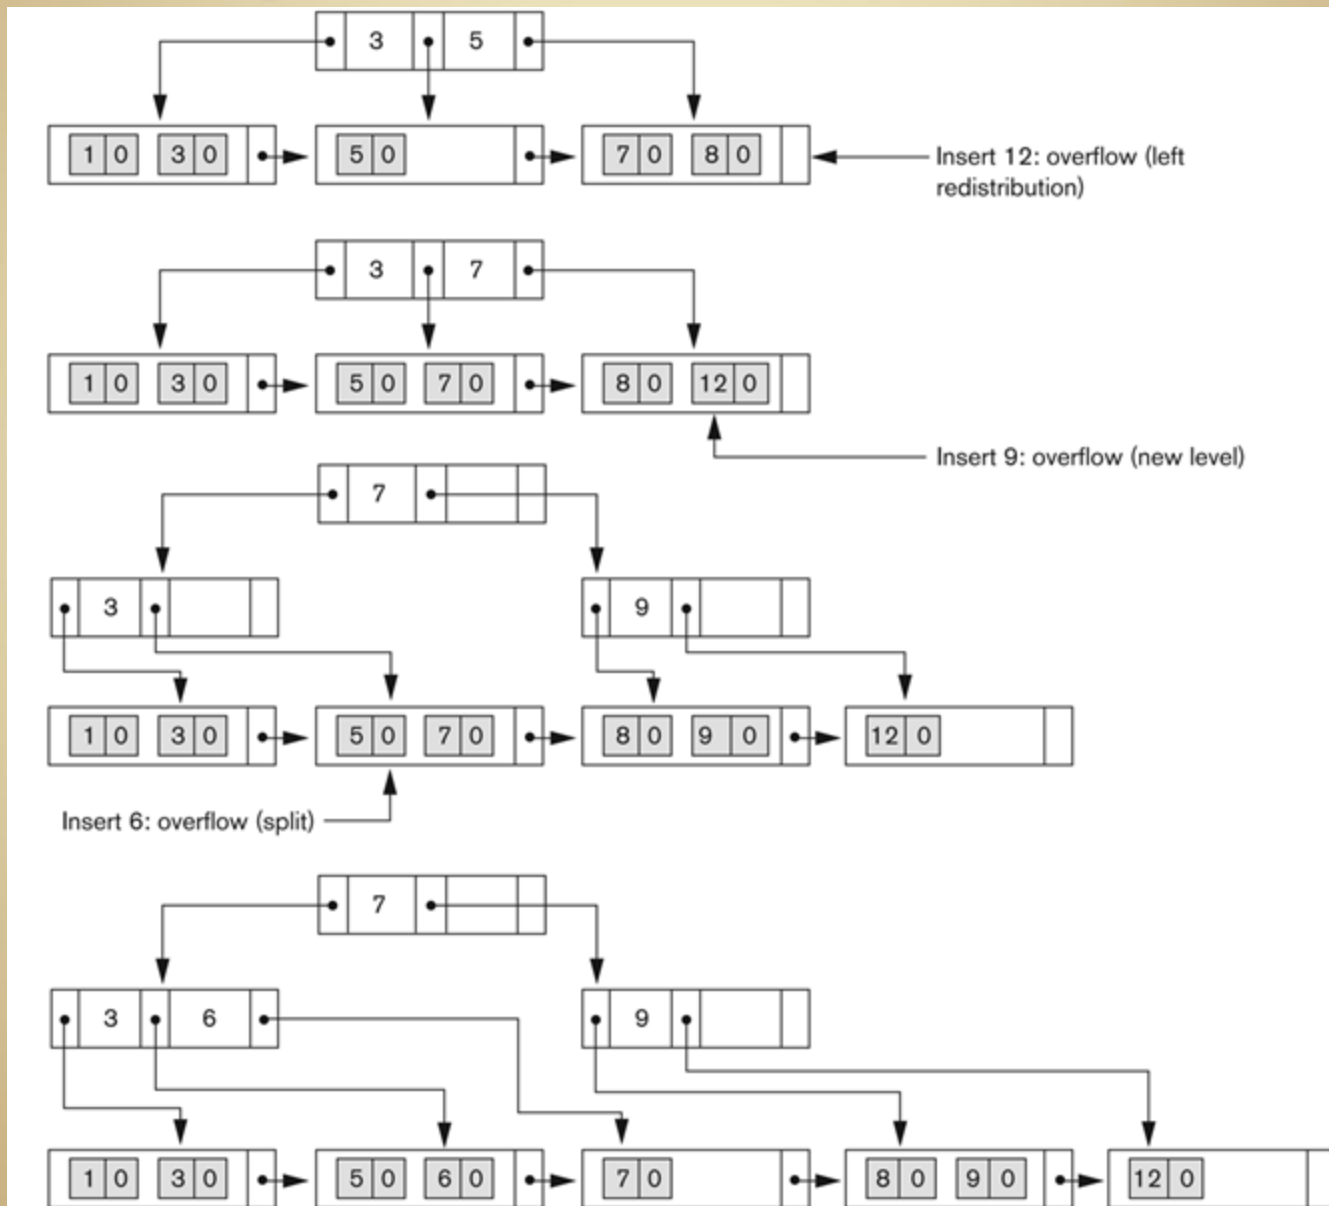
Example: B⁺-tree Insertion



Example: B⁺-tree Insertion



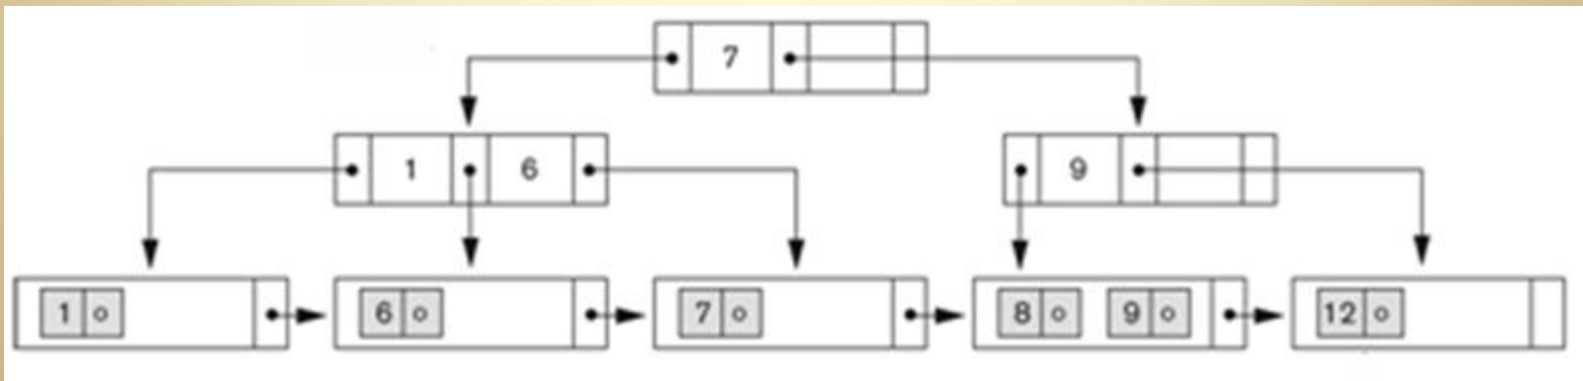
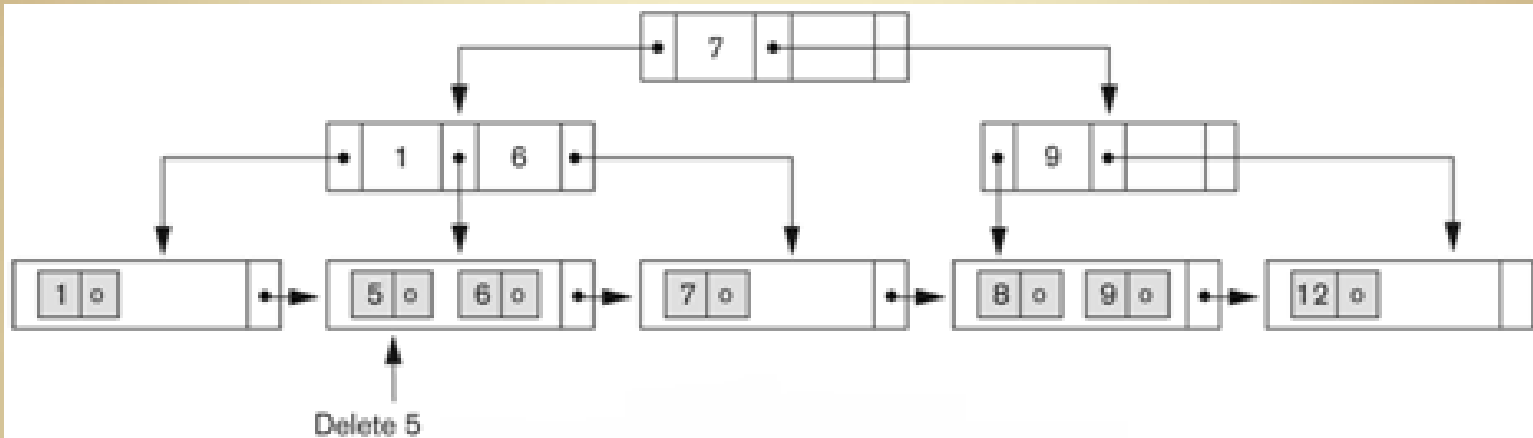
Example: B⁺-tree Insertion




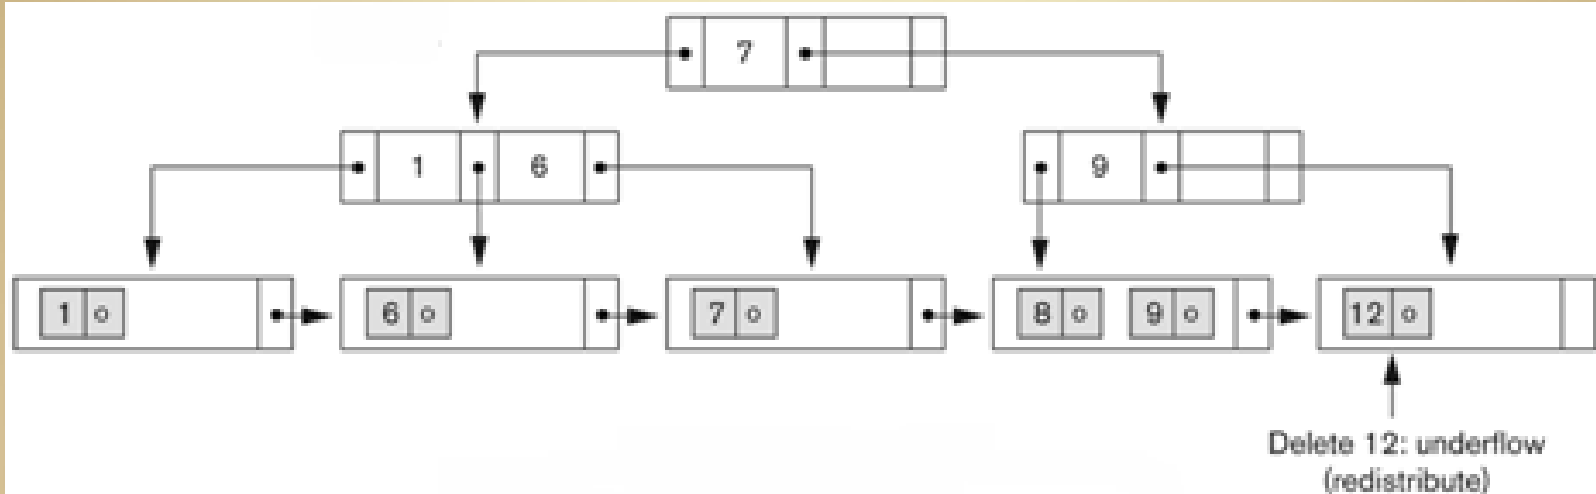
Handling Deletions

- Deletion can cause page to have fewer than q entries
 - Entries can be redistributed over adjacent pages to maintain minimum occupancy requirement
 - Ultimately, adjacent pages must be merged, and if merge propagates up the tree, height might be reduced
- In practice, tables generally grow, and merge algorithm is often not implemented
 - Reconstruct tree to compact it


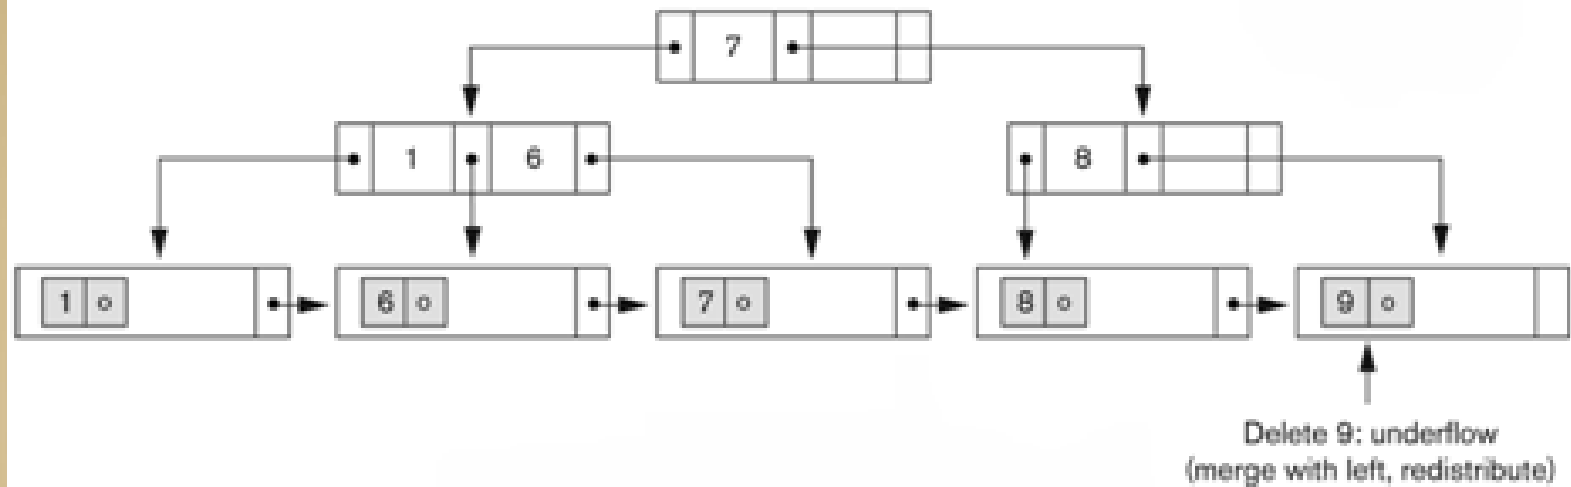
Example: B⁺-tree Deletion



Example: B⁺-tree Deletion



Example: B⁺-tree Deletion



REFERENCES

- Figures on slides 4, 5, 6, 20, 22, 25-28 borrowed from *Database Systems: An Application-Oriented Approach* (2nd ed) by Kifer, Bernstein and Lewis, Addison-Wesley, 2005
- Remaining figures from Elmasri/Navathe