

# **Advanced Database Management Systems**

**Lecture 18**  
**Query Processing: Select and Join**  
**Chapter 15**

# Algorithms for SELECT

- Consider only single table queries
- Three categories:
  - simple SELECT: one condition, no AND or OR
  - conjunctive select: multiple conditions, connected by AND
  - disjunctive select: multiple conditions, connected by OR

# Simple SELECT

- **Linear search (brute force):**
  - algorithm:
    - Retrieve every record in the file
    - test whether its attribute values satisfy the selection condition
  - works when:
    - always works
    - best on small files
    - only choice when no indexes or ordering
  - cost
    - average case:  $b/2$ , where  $b$  = # blocks in file
    - worst case:  $b$

# Simple SELECT

- **Binary search:**

- algorithm:
  - use binary search to find record(s)
- works when:
  - selection condition is an equality test on an ordering attribute
- cost:
  - $\log_2 b + \left\lceil \frac{s}{bfr} \right\rceil - 1$  , where  $b$  = # blocks in file,  $s$  = # selected records

# Simple SELECT

- **Primary index to retrieve a single record:**
  - algorithm:
    - look up record using primary index
  - works when:
    - selection condition is equality test on key attribute with primary index
  - cost:
    - $x + 1$ , where  $x = \#$  index levels

# Simple SELECT

- **Primary index to retrieve multiple records:**
  - algorithm:
    - use the index to find first record satisfying the corresponding equality condition
    - retrieve all matching subsequent records in the (ordered) file
  - works when:
    - selection condition is  $>$ ,  $\geq$ ,  $<$ , or  $\leq$  on key field with primary index
  - cost:
    - $x + \left\lceil \frac{s}{bfr} \right\rceil$ , where  $x$  = # index levels,  $s$  = # selected records

# Simple SELECT

- **Clustering index to retrieve multiple records:**
  - algorithm:
    - use clustering index to retrieve all the records satisfying the selection condition
  - works when:
    - selection condition is equality comparison on a non-key attribute with clustering index
  - cost:
    - $x + \left\lceil \frac{s}{bfr} \right\rceil$ , where  $x$  = # index levels,  $s$  = # selected records

# Simple SELECT

- **Secondary (B<sup>+</sup>-tree) index:**
  - algorithm:
    - lookup first record through B<sup>+</sup>-tree
    - scan leaves of B<sup>+</sup>-tree for additional records
  - works when:
    - selection is an equality test or range query on attribute that has a secondary index
    - selection is a range query on attribute that has a secondary index
  - cost:
    - equality test:  $x + s$  (worst case),  $x + 1$  (key)
    - range query:  $x + b_o/2 + r/2$



# Conjunctive SELECT

- **Conjunctive selection:**

- algorithm:
  - use one of the simple SELECT algorithms to find records matching one condition
  - check those records for remaining conditions
- works when:
  - conjunctive selection in which one of the simple SELECT algorithms can be applied to one condition
- cost:
  - same as simple SELECT cost
- example:
  - `select * from EMPLOYEE where DNO=6 and salary>70000,`  
and an index exists on DNO

# Conjunctive SELECT

- **Conjunctive selection using a composite index**
  - algorithm:
    - use composite index directly
  - works when:
    - selection condition is equality tests on two or more attributes for which a composite index exists
  - cost:
    - $x + 1$
  - example:
    - select \* from EMPLOYEE where LNAME='Jones' and FNAME='Sam', and an index exists on <LNAME, FNAME>

# Conjunctive SELECT

- **Conjunctive selection by intersection of record pointers:**
  - algorithm:
    - use indexes to find record pointers for each equality condition
    - compute intersection of record pointer sets
    - retrieve records using records pointers
  - works when:
    - secondary indexes are available on all (or some of) the fields involved in equality comparison conditions
    - indexes include record pointers (rather than block pointers)
  - cost:
    - sum of index operations plus  $s$ , where  $s$  = #records selected
  - example:
    - select \* from WORKS ON where PNO=10 and LNAME='Wolfe', and indexes exist for both PNO and LNAME

# Disjunctive SELECT

- Disjunctive selects are much harder to optimize
  - no single condition can be used to 'pre-filter' the results
  - result is union of each condition
  - best you can do is to try to optimize each individual query, then compute the union
  - example:
    - `select * from EMPLOYEE where DNO=3 or SALARY>80000 or SEX='F'`

# JOIN Algorithms

- We'll consider joins such as  $R \bowtie_{A=B} S$
- Extends to joins like  $R \bowtie_{A=B \text{ and } C=D} S$   
by considering  $\langle A, C \rangle$  and  $\langle B, D \rangle$  as single attributes

# Algorithms for JOIN: $R \bowtie_{A=B} S$

- Nested loop (brute force)
  - algorithm:

```
for r in R
  for s in S
    if (r[A] = s[B])
      write < r, s >
```
  - Cost:
    - both tables and result can fit in memory:  
 $b_R + b_S$ , where  $b_R = \text{\#blocks in R}$  and  $b_S = \text{\#blocks in S}$
    - both tables and result cannot fit in memory:  
 $b_R * b_S$ , where  $b_R = \text{\#blocks in R}$  and  $b_S = \text{\#blocks in S}$

# Algorithms for JOIN: $R \bowtie_{A=B} S$

- Single loop
  - algorithm: use an access structure on S:  
    for r in R  
        for s in  $\sigma_{r[A]=s[B]} S$   
            write  $\langle r, s \rangle$
  - Cost:
    - $b_R * (x+1)$ , where  $b_R$  = # blocks in R and x = levels in index on S

# Algorithms for JOIN: $R \bowtie_{A=B} S$

- Sort-Merge

- works when R is ordered on A and S is ordered on B
- algorithm:
  - Scan both files in order of the join attributes, matching the records that have the same values for A and B
- Cost:
  - $b_R + b_S$ , where  $b_R = \text{\#blocks in R}$  and  $b_S = \text{\#blocks in S}$



# Algorithms for JOIN: $R \bowtie_{A=B} S$

- Hash-Join

- algorithm:

- assume R is smaller table (without loss of generality)
    - Scan through smaller file, R, and hash records on A
    - Scan through other file, S, and hash records on B (using same hash function)
    - As S's records are hashed:  
if matching record from R is in the bucket, build and store the result tuple, otherwise

- Cost:

- $b_R + b_S$ , where  $b_R = \text{\#blocks in R}$  and  $b_S = \text{\#blocks in S}$
    - works when R can fit in memory

# Algorithms for JOIN: $R \bowtie_{A=B} S$

- **Hash-join:**

- The records of files R and S are both hashed to the *same hash file*, using the *same hashing function* on the join attributes A of R and B of S as hash keys.
- A single pass through the file with fewer records (say, R) hashes its records to the hash file buckets.
- A single pass through the other file (S) then hashes each of its records to the appropriate bucket, where the record is combined with all matching records from R.

# JOIN Cost Functions

nested-loop	$b_R + (b_R * b_S) + ((js *  R  *  S ) / bfr_{RS})$
single-loop	using secondary index $b_R + ( R  * (x_B + s_B)) + ((js *  R  *  S ) / bfr_{RS})$ using primary index $b_R + ( R  * (x_B + 1)) + ((js *  R  *  S ) / bfr_{RS})$
sort-merge	$b_R + b_S + ((js *  R  *  S ) / bfr_{RS})$

js: join selectivity = ratio of join size vs. cross-product size

$(js * |R| * |S|) / bfr_{RS}$  : cost to write result back to disk