

L e c t u r e



30

Review of Last Lecture

- RFCs
- HTTP protocol
- MIME types
- Server Architecture
- Threads

Server Architecture

- Dialog-based GUI application
- Most of the processing is at back-end
- Running on TCP port 5432 decimal

HTTP Web Server Application

Initialise Windows Sockets Library

```
if (WSAStartup(MAKEWORD(1,1), &wsaData))  
{  
    ... ..  
    return 1;  
}
```

HTTP Web Server Application

Get machine's hostname and IP address

```
gethostname(hostName, sizeof(hostName));  
ptrHostEnt = gethostbyname(hostName);
```

Fill the socket address with appropriate values

```
serverSocketAddress.sin_family = AF_INET;  
serverSocketAddress.sin_port = htons(SERVER_PORT);  
... ..  
memcpy(&serverSocketAddress.sin_addr.S_un.S_addr,  
ptrHostEnt->h_addr_list[0], sizeof(unsigned long));
```

HTTP Web Server Application

Create the server socket

```
serverSocket = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP);  
if(serverSocket == INVALID_SOCKET)  
{  
    ... ..  
    WSACleanup();  
    return 1;  
}
```

Bind the socket

```
if(bind(serverSocket,  
        (struct sockaddr *)&serverSocketAddress,  
        sizeof(serverSocketAddress)))  
{  
    ... ..  
    WSACleanup();  
    return 1;  
}
```

HTTP Web Server Application

Put the socket in listening mode

```
if(listen(serverSocket, MAX_PENDING_CONNECTIONS))  
{  
    ... ..  
    WSACleanup();  
    return 1;  
}
```

Here is the time to accept client connections

HTTP Web Server Application

Create a thread that will call `accept()` in a loop to accept multiple client connections

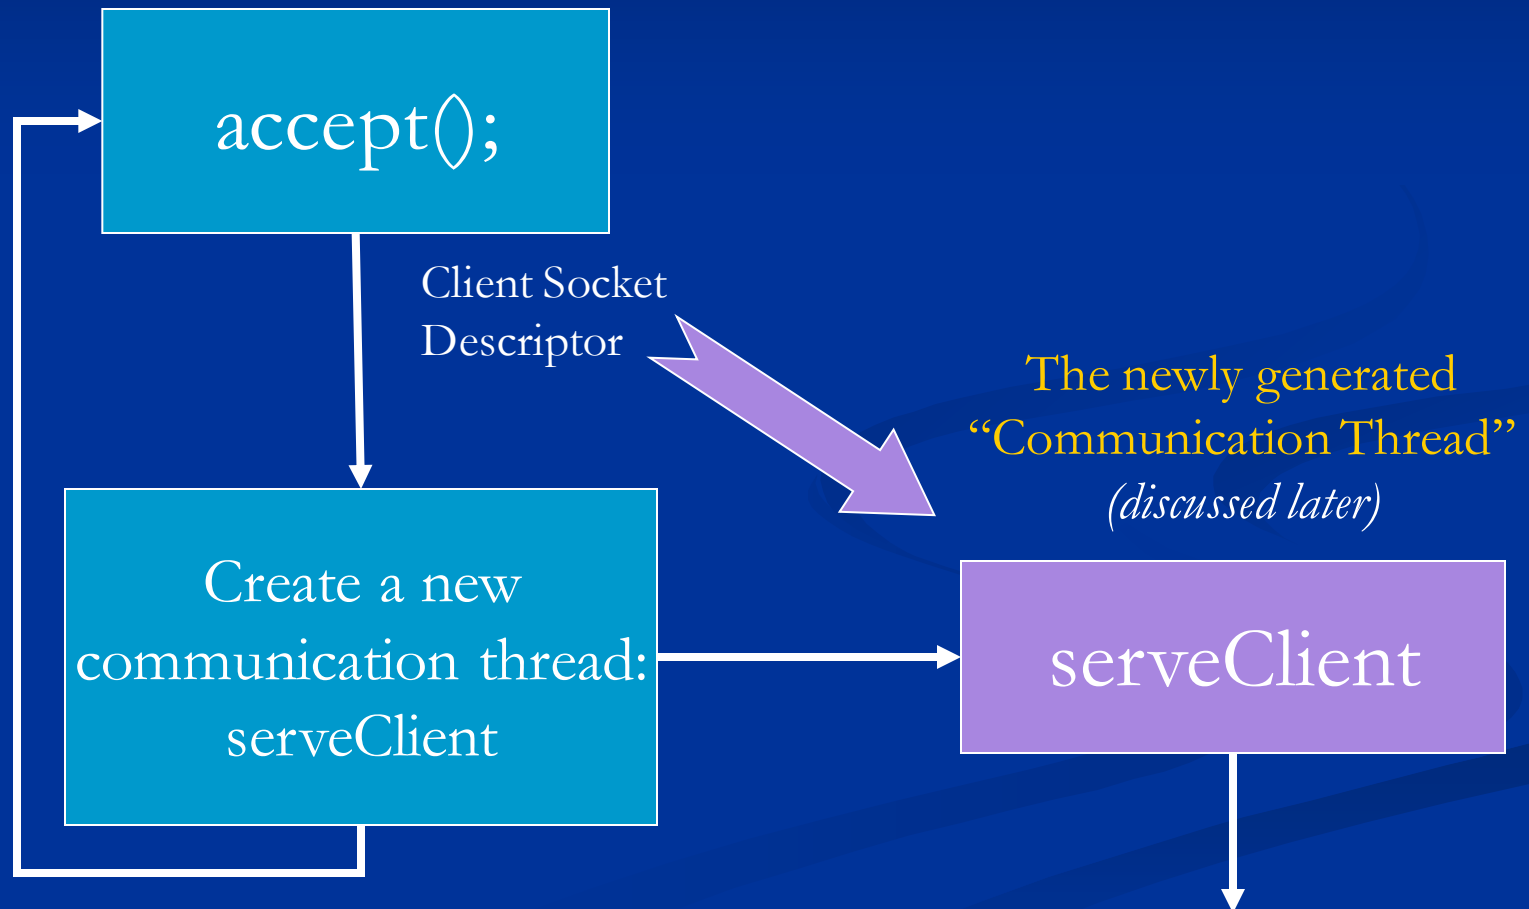
```
hAcceptingThread = CreateThread(  
    NULL,  
    0,  
    (LPTHREAD_START_ROUTINE)  
    acceptClientConnections,  
    NULL,  
    CREATE_SUSPENDED,  
    &dwAcceptingThread);
```


HTTP Web Server Application

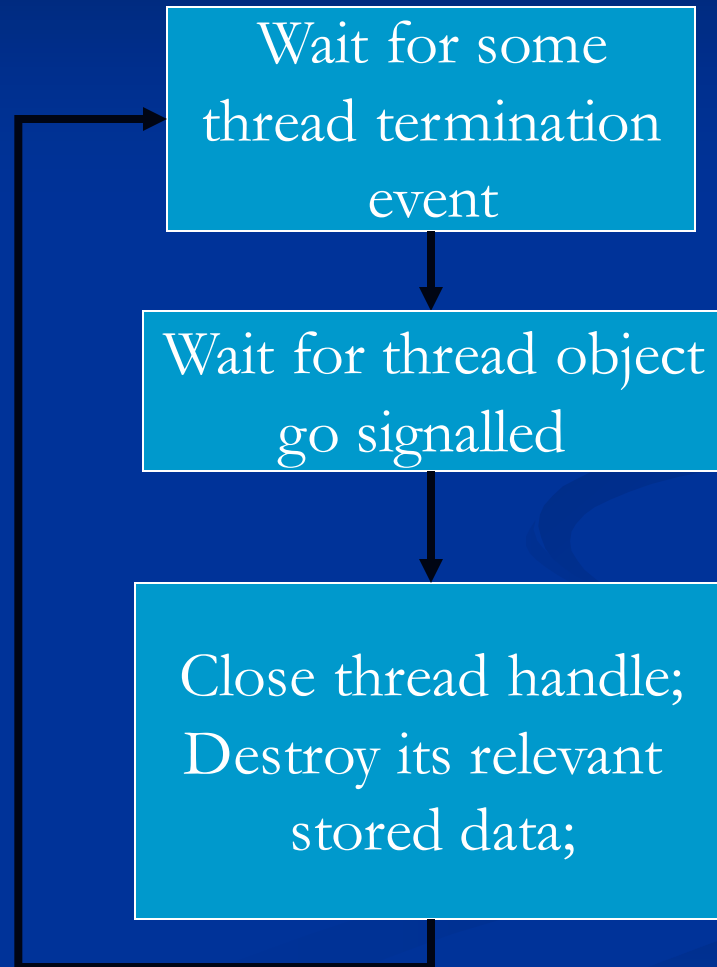
Create a thread to do termination house-keeping
when some communication thread terminates

```
hTerminatingThread = CreateThread(  
    NULL,  
    0,  
    (LPTHREAD_START_ROUTINE)  
    terminateCommunicationThreads,  
    NULL,  
    CREATE_SUSPENDED,  
    &dwTerminatingThread);
```

acceptClientConnections thread routine



terminateCommunicationThreads thread routine



Application Variables

```
#define MAX_CLIENTS 5
```

```
SOCKET clientSockets[MAX_CLIENTS];
```

```
HANDLE hCommunicationThreads[MAX_CLIENTS];
```

```
DWORD dwCommunicationThreads[MAX_CLIENTS];
```


```
HANDLE hAcceptingThread;
```

```
DWORD dwAcceptingThread;
```

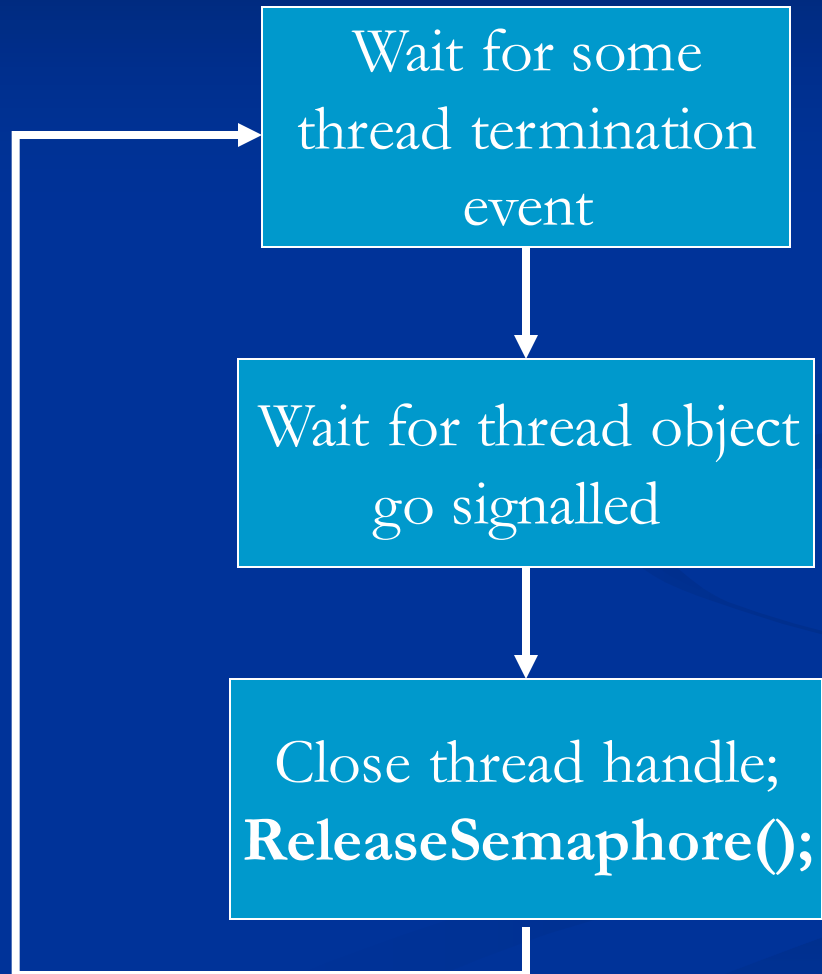
```
HANDLE hTerminatingThread;
```

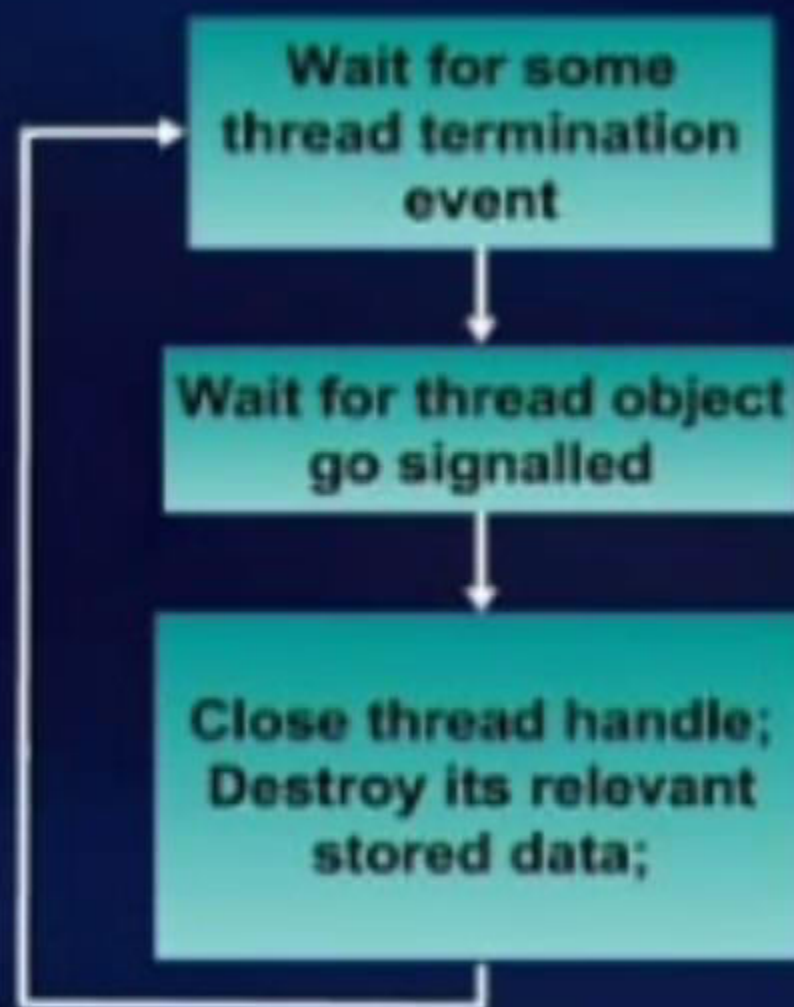
```
DWORD dwTerminatingThread;
```

acceptClientConnections thread routine



terminateCommunicationThreads thread routine





serveClient Communication Thread routine

Communicate with client to receive/serve its HTTP request
Use `recv()` / `send()` blocking WinSock API calls

HTTP request served
going to disconnect the client

Set an **Event** object to indicate termination

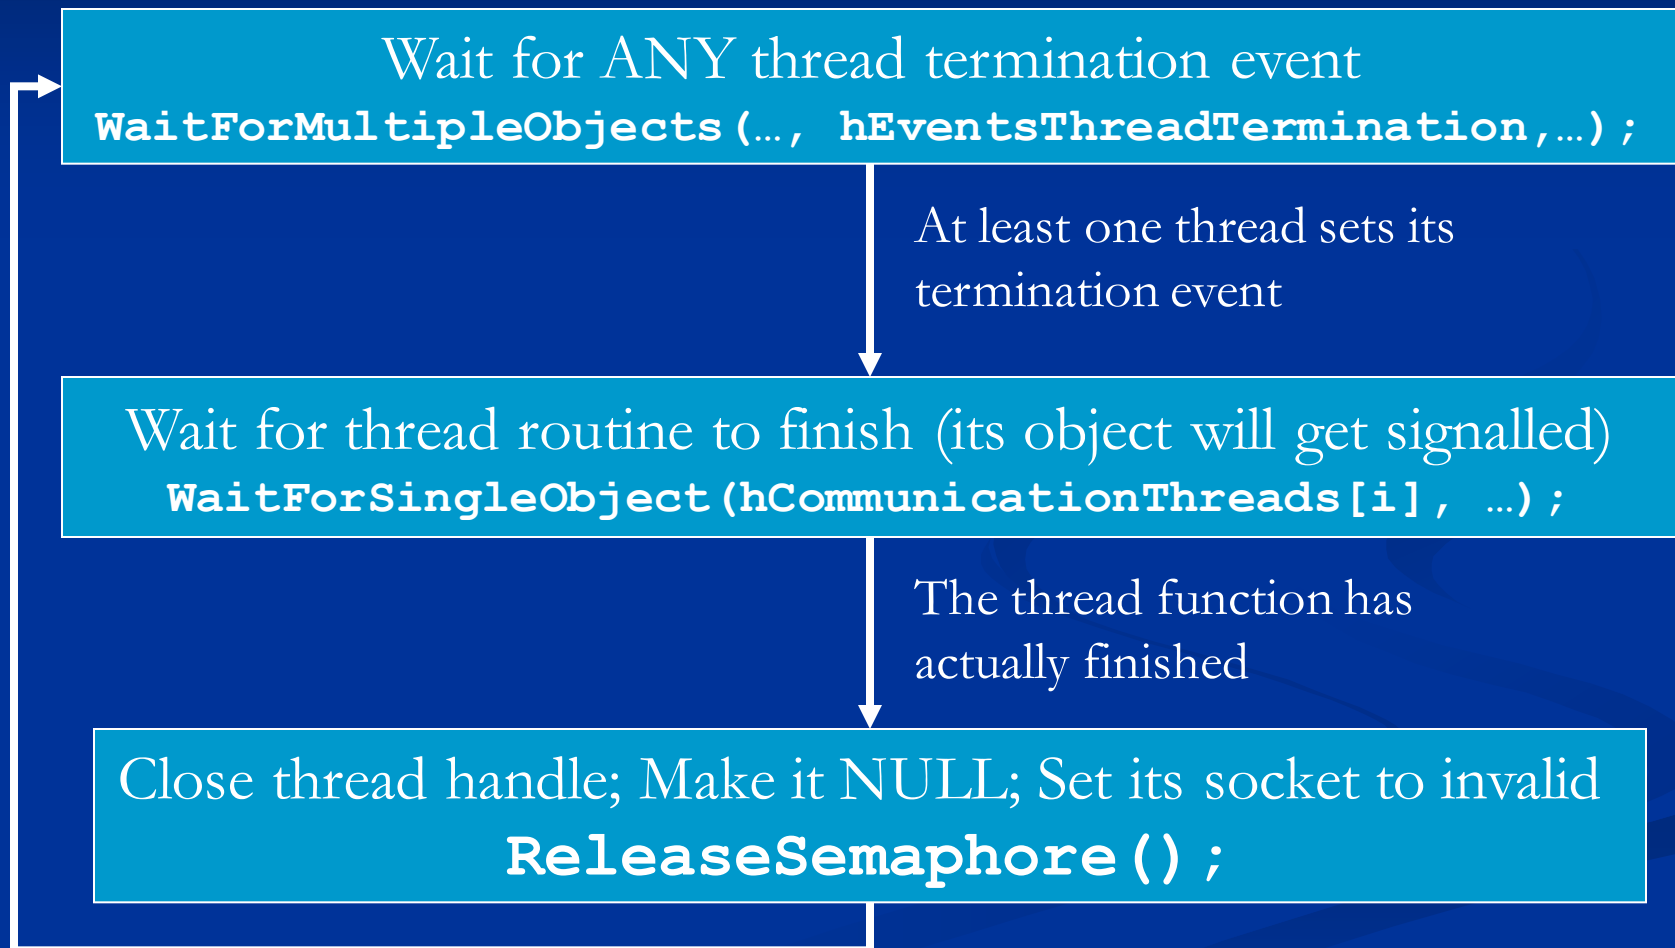
Gracefully shutdown and Close client socket

Application Variables

```
#define MAX_CLIENTS    5
```

```
HANDLE hEventsThreadTermination[MAX_CLIENTS];
```

terminateCommunicationThreads thread routine



Thread Procedures Summary

acceptClientConnections

- to accept client connection

terminateCommunicationThreads

- to do housekeeping when communication threads terminate

serveClient

- to do actual communication to receive and serve an HTTP request

User Interface



Variable Initialisation

```
for(i=0; i<MAX_CLIENTS; ++i)
{
    clientSockets[i] = INVALID_SOCKET;

    hCommunicationThreads[i] = NULL;
    dwCommunicationThreads[i] = 0;

    hEventsThreadTermination[i] = NULL;
}
```

Initialise WinSock Library

```
if (WSAStartup (MAKEWORD (1,1) , &wsaData))  
{  
    MessageBox (NULL,  
        "Error initialising sockets library.",  
        "WinSock Error",  
        MB_OK | MB_ICONSTOP) ;  
  
    return 1;  
}
```

Win32 Error Codes

```
int  WSAGetLastError(void) ;
```

- get error code for the last unsuccessful Windows Sockets operation

```
DWORD GetLastError(VOID) ;
```

- retrieve calling threads last-error code

HTTP Web Server Application

Get machine's hostname and IP address

```
gethostname(hostName, sizeof(hostName));  
ptrHostEnt = gethostbyname(hostName);
```

Fill the socket address with appropriate values

```
serverSocketAddress.sin_family = AF_INET;  
serverSocketAddress.sin_port = htons(SERVER_PORT);  
... ..  
memcpy(&serverSocketAddress.sin_addr.S_un.S_addr,  
ptrHostEnt->h_addr_list[0], sizeof(unsigned long));
```


HTTP Web Server Application

Create the server socket

```
serverSocket = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP);  
if(serverSocket == INVALID_SOCKET)  
{  
    ... ..  
    WSACleanup();  
    return 1;  
}
```

Bind the socket

```
if(bind(serverSocket,  
        (struct sockaddr *)&serverSocketAddress,  
        sizeof(serverSocketAddress)))  
{  
    ... ..  
    WSACleanup();  
    return 1;  
}
```

HTTP Web Server Application

Put the socket in listening mode

```
if(listen(serverSocket, MAX_PENDING_CONNECTIONS))  
{  
    ... ..  
    WSACleanup();  
    return 1;  
}
```

Here is the time to accept client connections

Limiting Maximum Concurrent Connections

Create an unnamed semaphore object with
MAX_CLIENTS as initial/maximum count

```
hSemaphoreMaxClients =  
    CreateSemaphore(NULL,  
                    MAX_CLIENTS,  
                    MAX_CLIENTS,  
                    NULL) ;
```

“I am dying...”, the thread said

create an array of non-signalled event
objects

```
for (i=0; i<MAX_CLIENTS; i++)  
    hEventsThreadTermination[i] =  
        CreateEvent(NULL, FALSE, FALSE, NULL);
```

HTTP Web Server Application

Create the connection-accepting thread

```
hAcceptingThread = CreateThread(  
    NULL,  
    0,  
    (LPTHREAD_START_ROUTINE)  
        acceptClientConnections,  
    NULL,  
    CREATE_SUSPENDED,  
    &dwAcceptingThread);
```

Create the termination house-keeping thread

```
hTerminatingThread = CreateThread(... ..);
```

Display the dialog

```
DialogBox(..., ..., ..., mainDialogProc);
```

mainDialogProc dialog procedure

Resume the threads

```
case WM_INITDIALOG:  
    ResumeThread(hAcceptingThread) ;  
    ResumeThread(hTerminatingThread) ;  
    return TRUE ;  
    break ;
```

Handling the server shut-down button

```
case IDC_BUTTON_SHUTDOWN:  
    Perform any shut-down tasks that may be necessary  
    EndDialog(hDlg, 0) ;  
    break ;
```

acceptClientConnections thread routine

Start of the loop to accept client connections

Wait for semaphore count to go non-zero

```
dwWaitResult = WaitForSingleObject(  
    hSemaphoreMaxClients,  
    INFINITE) ;
```

```
switch(dwWaitResult)  
{  
case WAIT_OBJECT_0:
```

We can accept more connections here because
semaphore object is signalled

```
clientSocket = accept(... ..) ;
```

acceptClientConnections thread routine

```
clientSocket = accept(... ..) ;
```

Connection accepted! Look for the first empty slot to save the new socket descriptor

```
for(i=0; i<MAX_CLIENTS; i++)  
{  
    if(clientSockets[i] == INVALID_SOCKET)  
        break;  
}  
  
nextClientIndex = i;  
  
clientSockets[nextClientIndex]=clientSocket;
```

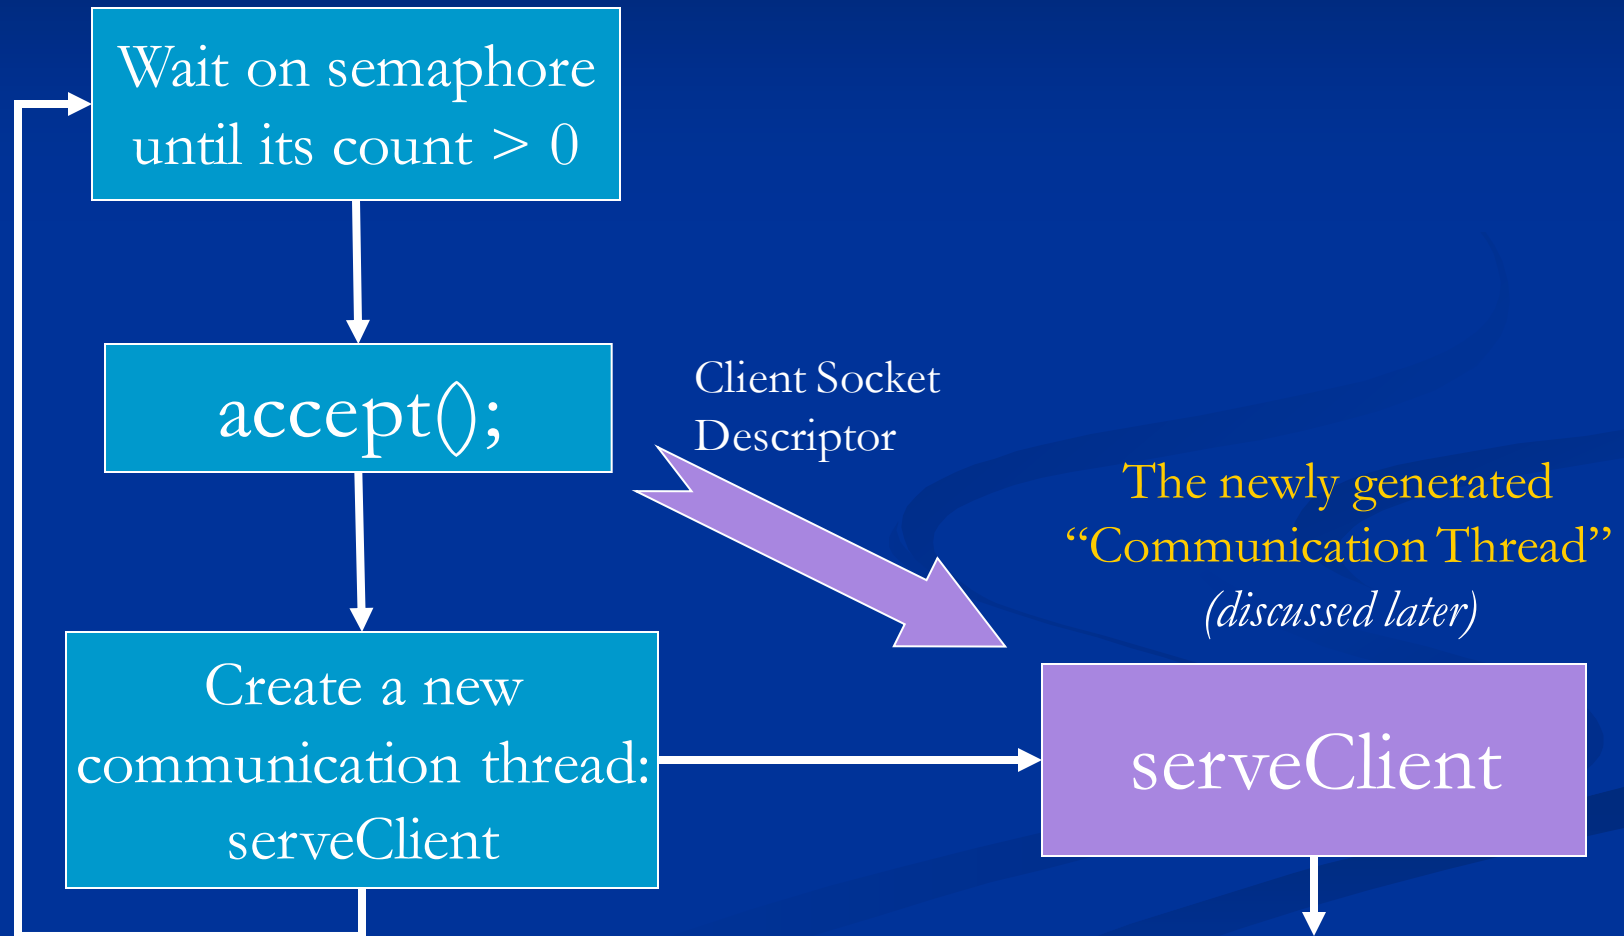

acceptClientConnections thread routine

`nextClientIndex` is used as an index in ALL arrays to store information relevant to this new client connection

```
clientSockets[nextClientIndex]=clientSocket;
```

```
hCommunicationThreads[nextClientIndex] =  
    CreateThread(..., ...,  
        serveClient,                thread procedure  
        (LPVOID)nextClientIndex,    thread parameter  
        CREATE_SUSPENDED,  
        ...);
```

acceptClientConnections thread routine



serveClient thread routine

Index for this client in all arrays is passed to this thread routine

```
DWORD WINAPI serveClient(LPVOID clientNumber)
{
    char msg[2046] = "";
```

Receiving an HTTP request from browser

```
recv(
    clientSockets[(UINT)clientNumber],
    msg,
    2046,
    0 );
```

Available HTTP request data
will be received in this buffer

acceptClientConnections thread routine

`nextClientIndex` is used as an index in ALL arrays to store information relevant to this new client connection

```
clientSockets[nextClientIndex]=clientSocket;
```

```
hCommunicationThreads[nextClientIndex] =  
    CreateThread(..., ...,  
        serveClient,
```

```
        (LPVOID)nextClientIndex, thread parameter  
        CREATE_SUSPENDED,  
        ...);
```

Sample Request

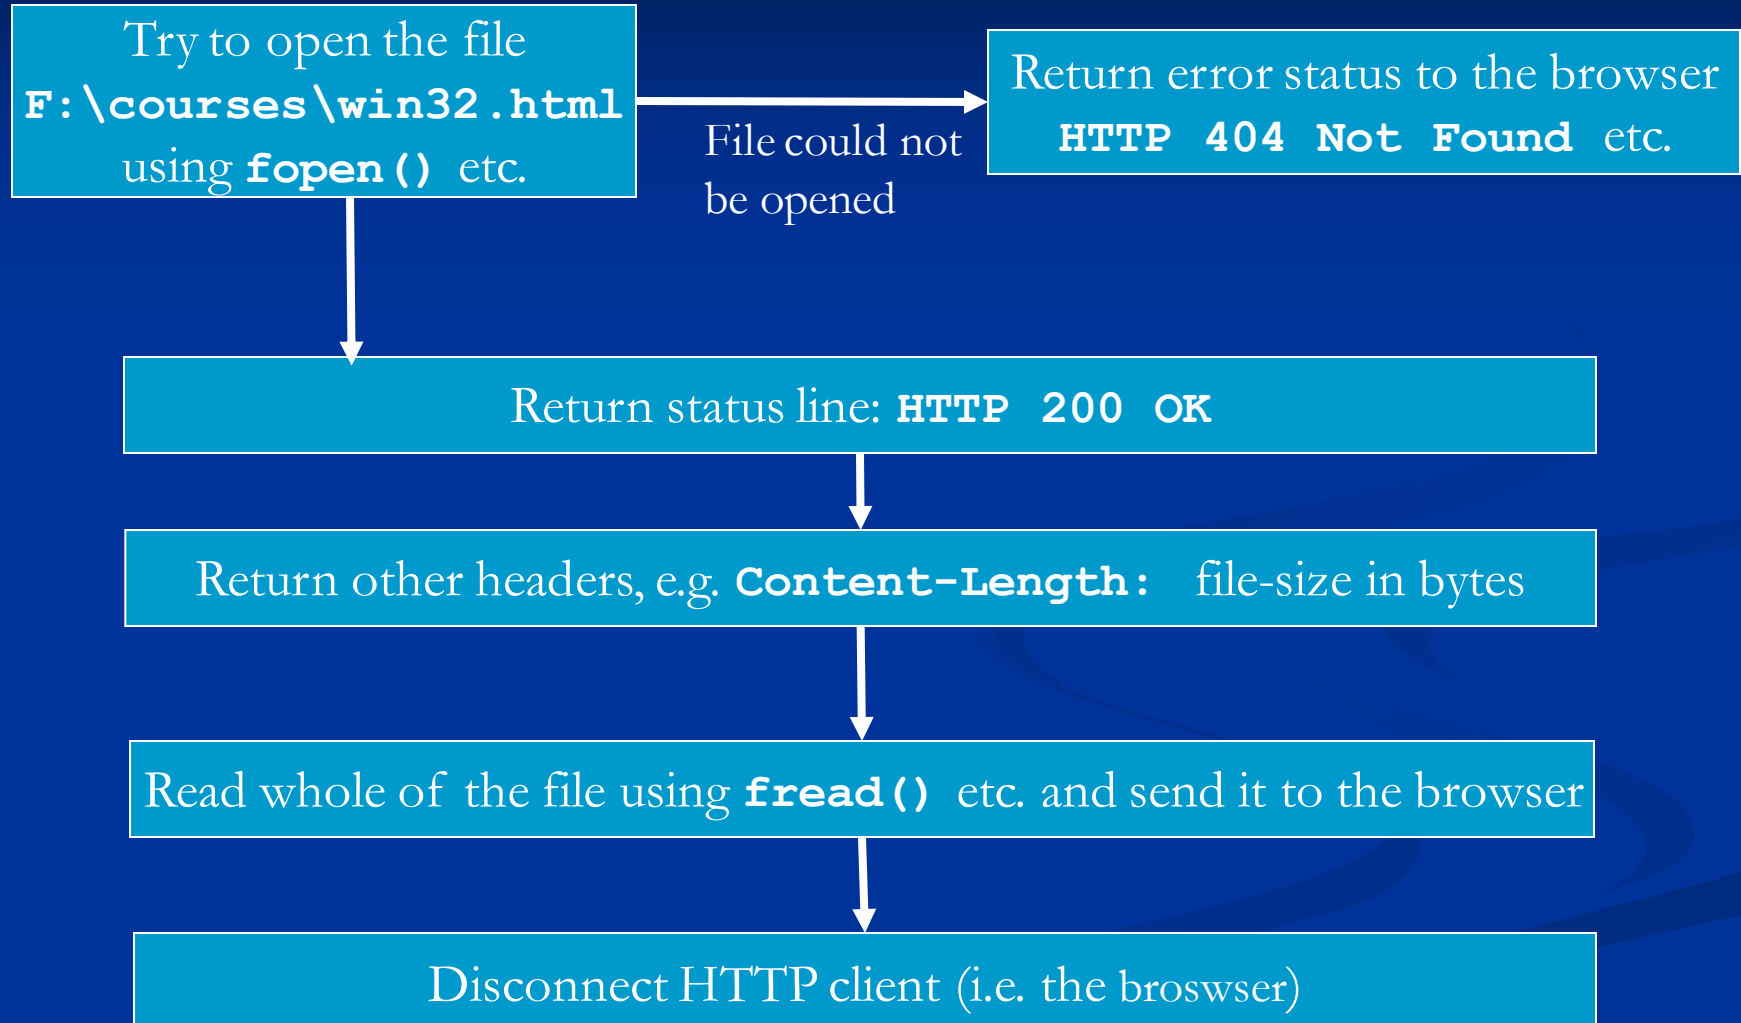
- Request parsing: understanding what the client has demanded

GET **/courses/win32.html** HTTP/1.0

- Assume **F:** is your server's home directory, and **\courses** is not a virtual directory, server should return the file

F:\courses\win32.html

Server returns the requested file



HTTP Redirection

- Redirecting the client irrespective of the HTTP request!

The string in the #define directive is assumed to be on a single line

```
#define RESPONSE
```

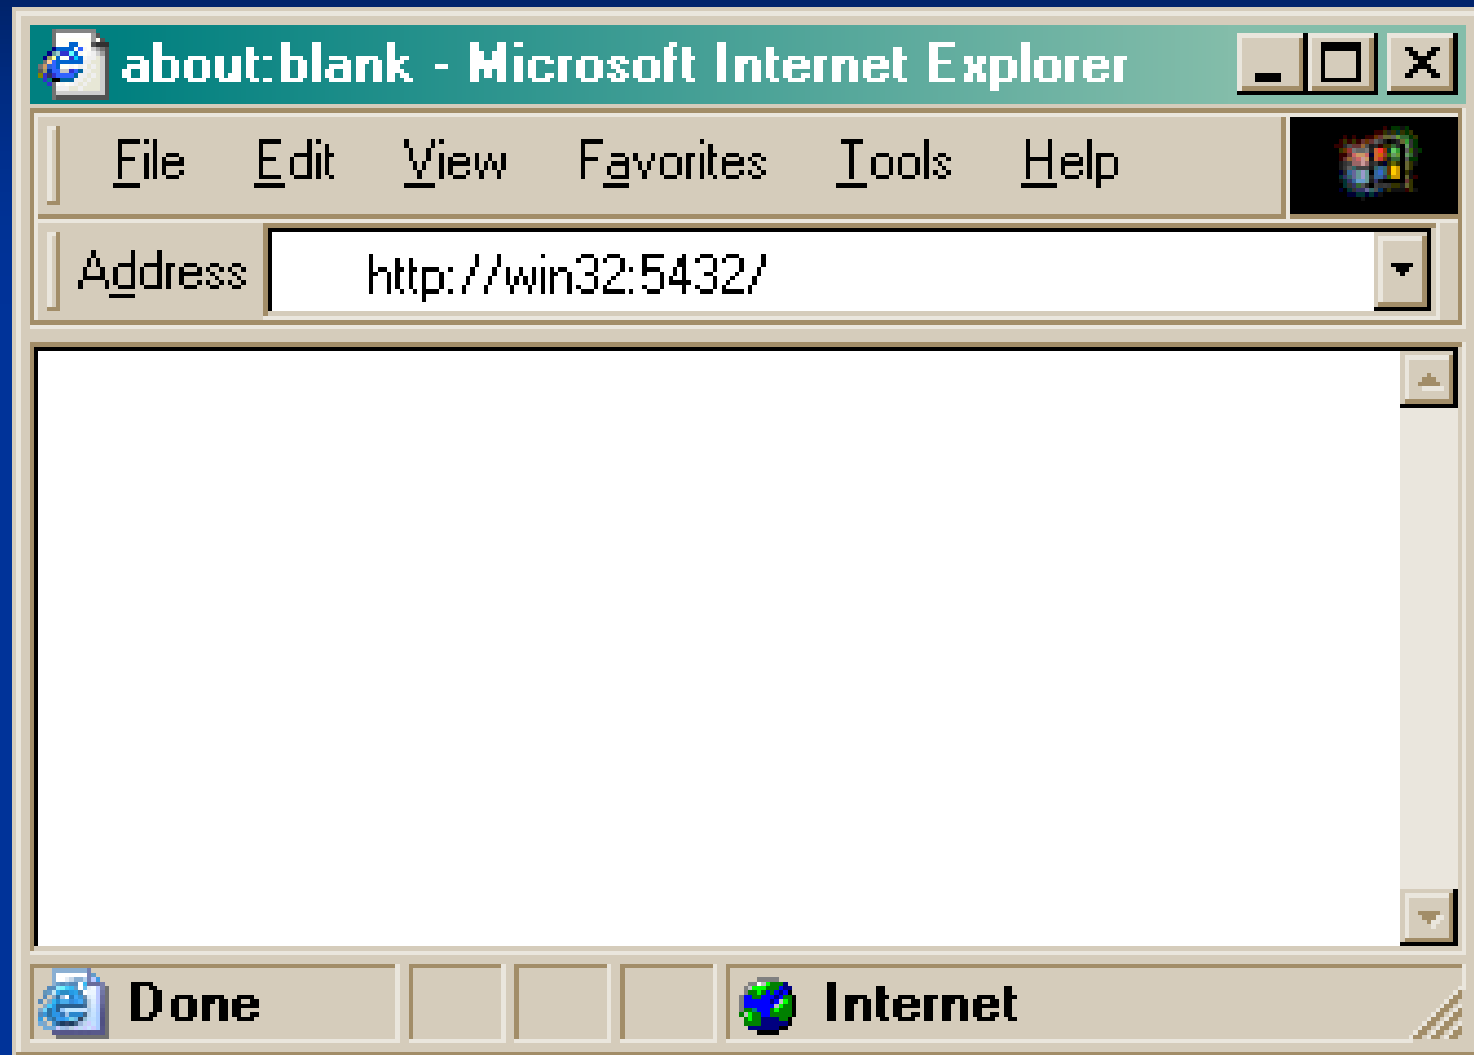
```
"HTTP/1.1 302 Object Moved\r\n  
Location: http://www.vu.edu.pk\r\n\r\n"
```

Status line

Sending the hard-coded HTTP response back to browser

```
send(clientSockets[ (UINT)clientNumber] ,  
RESPONSE ,  
sizeof (RESPONSE) ,  
0) ;
```

Sending a request to our Web Server



Using Port Numbers

- There is no compulsion to build all HTTP Web Servers to run on port 80
- These are ‘suggested’ port numbers for a Win32 developer
- Standard servers do run on port 80
- Our HTTP Web Server may also need to run on port 80 if put it to public use

Returning an HTML document

The string in the #define directive is assumed to be on a single line

```
#define RESPONSE
```

```
"HTTP/1.0 200 OK\r\n  
Content-type: text/html\r\n  
Content-length: 1325\r\n\r\n"
```

File **a.htm** is
1325 bytes long

Send the hard-coded HTTP status and headers

```
send(clientSockets[(UINT)clientNumber], RESPONSE,  
      sizeof(RESPONSE), 0);
```

Now send the whole file using character I/O of standard C runtime

```
ch = fgetc(fptr);  
while(!feof(fptr)) {  
    send(clientSockets[(UINT)clientNumber], &ch, 1, 0);  
    ch = fgetc(fptr);  
}
```

A Flawed Web Server

- Fixed sized arrays waste memory and lack run-time flexibility
- One event per thread to signify termination:
`WaitForMultipleObjects()` can not wait on more than a certain number of objects e.g. 64 on x86 under NT.

Dynamic Web Content

- Server blindly dumps HTML files to the clients.
This is 'static content'.
- If server reads the file and modifies its output
e.g.
`%%time%%` replaced with current system time
Every 2 clients connected at different instants of
time will receive **different** content.
This is 'dynamic content'.
- `%%time%%` may be called a **tag**

Dynamic Web Content

- Microsoft Active Server Pages
- Macromedia ColdFusion
- Tags are not sent to the client. These are processed by the server and the resulting output is sent to the browser.

CGI

- Common Gateway Interface
- Win32 EXE executable is executed at the server
- All browser request data is available at **stdin** (read using **scanf()** etc.)
- All output sent to **stdout** (output using **printf** etc.) is sent to the browser instead of the server screen.